

Title of the dissertation

A DISSERTATION PRESENTED

BY

FIRSTNAME M. LASTNAME

TO

THE DEPARTMENT OF PSYCHOLOGY

IN PARTIAL FULFILLMENT OF THE REQUIREMENTS

FOR THE DEGREE WITH HONORS OF

DOCTOR OF PHILOSOPHY

IN THE SUBJECT OF

PSYCHOLOGY

HARVARD UNIVERSITY

CAMBRIDGE, MASSACHUSETTS

MAY 2024

©2014 – FIRSTNAME M. LASTNAME
ALL RIGHTS RESERVED.

Title of the dissertation

ABSTRACT

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Morbi commodo, ipsum sed pharetra gravida, orci magna rhoncus neque, id pulvinar odio lorem non turpis. Nullam sit amet enim. Suspendisse id velit vitae ligula volutpat condimentum. Aliquam erat volutpat. Sed quis velit. Nulla facilisi. Nulla libero. Vivamus pharetra posuere sapien. Nam consectetur. Sed aliquam, nunc eget euismod ullamcorper, lectus nunc ullamcorper orci, fermentum bibendum enim nibh eget ipsum. Donec porttitor ligula eu dolor. Maecenas vitae nulla consequat libero cursus venenatis. Nam magna enim, accumsan eu, blandit sed, blandit a, eros.

Quisque facilisis erat a dui. Nam malesuada ornare dolor. Cras gravida, diam sit amet rhoncus ornare, erat elit consectetur erat, id egestas pede nibh eget odio. Proin tincidunt, velit vel porta elementum, magna diam molestie sapien, non aliquet massa pede eu diam. Aliquam iaculis. Fusce et ipsum et nulla tristique facilisis. Donec eget sem sit amet ligula viverra gravida. Etiam vehicula urna vel turpis. Suspendisse sagittis ante a urna. Morbi a est quis orci consequat rutrum. Nullam egestas feugiat felis. Integer adipiscing semper ligula. Nunc molestie, nisl sit amet cursus convallis, sapien lectus pretium metus, vitae pretium enim wisi id lectus. Donec vestibulum. Etiam vel nibh. Nulla facilisi. Mauris pharetra. Donec augue. Fusce ultrices, neque id dignissim ultrices, tellus mauris dictum elit, vel lacinia enim metus eu nunc.

Contents

1	INTRODUCTION	I
2	BACKGROUND	2
2.1	Large Language Models	2
2.2	Hallucination in Large Language Models	4
2.2.1	Taxonomy	4
2.2.2	Causes	5
2.2.3	Consequences and Detection Challenges	6
2.3	Text Embeddings and Embedding Geometry	6
2.3.1	Text Embeddings	7
2.3.2	Cosine Distance	7
2.3.3	High-Dimensional Geometry	7
2.3.4	The Manifold Hypothesis	8
2.3.5	Geometric Concepts	8
2.4	Evaluating Language Model Outputs	10
2.4.1	The Evaluation Bottleneck	10
2.4.2	LLM-as-Judge	10
2.4.3	Known Biases	11
2.4.4	Rubric Design	11
2.5	Fine-Tuning and Parameter-Efficient Adaptation	12
2.5.1	Supervised Fine-Tuning	12
2.5.2	LoRA: Low-Rank Adaptation	12
2.5.3	Self-Distillation	13
3	LITERATURE REVIEW	14
3.1	Hallucination: Taxonomy and Theoretical Limits	14
3.2	Benchmarks for Factuality Evaluation	15
3.3	Automated Evaluation and the LLM-as-Judge Paradigm	16
3.4	Detecting Hallucination from Internal Representations	17
3.4.1	Probing for Truthfulness	17
3.4.2	Geometric and Spectral Methods	17
3.4.3	Uncertainty-Based and Inference-Time Methods	18
3.5	Entity Knowledge and Knowledge Boundaries	19
3.6	Prompt Engineering for Factuality	20

3.7	Fine-Tuning for Truthfulness	21
3.7.1	Alignment and Its Costs	21
3.7.2	Factuality-Targeted Fine-Tuning	21
3.7.3	Data Selection	21
3.8	Synthesis and Research Questions	22
4	EXPERIMENTAL SETUP	24
4.1	Benchmark Construction	25
4.1.1	Category Design	25
4.1.2	Template Design	29
4.1.3	Entity Pools and Validation	30
4.1.4	Benchmark Scaling	32
4.2	Model Selection	34
4.2.1	Cross-Model Benchmark Suite	35
4.2.2	Intervention Target Models	36
4.2.3	Generation Parameters	38
4.3	Consensus Judging Pipeline	39
4.3.1	Why LLM Judges	39
4.3.2	Panel Design	40
4.3.3	Evaluation Rubric	41
4.3.4	Judge Prompt Design	43
4.3.5	Majority Vote and Consensus	45
4.3.6	Human Validation	46
4.3.7	Sensitivity Analysis and Limitations	47
4.4	Embedding and Geometric Feature Extraction	48
4.4.1	Embedding Model	48
4.4.2	Corpus Construction	49
4.4.3	Geometric Features	50
4.4.4	Hyperparameter Choices	55
4.4.5	Limitations	56
4.5	Summary of Experimental Infrastructure	57
4.6	Statistical Methods	58
5	CAN PROMPTS REDUCE HALLUCINATIONS?	60
5.1	Prefix Design	60
6	CAN GEOMETRY GUIDE INTERVENTION?	63
7	DISCUSSION	64
	APPENDIX A EXPERIMENTAL ARTIFACTS	65

A.1	Judge System Prompt	65
A.2	Prompt Prefix Texts	67
A.3	Fine-Tuning Configurations	69
APPENDIX B SUPPLEMENTARY ANALYSIS		71
B.1	Cross-Model Correlation Matrix	71
B.2	Embedding Model Robustness	71
B.3	Human Validation Details	71
B.4	Geometric Feature Visualizations	71
REFERENCES		78
REFERENCES		87

THIS IS THE DEDICATION.

Acknowledgments

LOREM IPSUM DOLOR SIT AMET, consectetur adipiscing elit. Morbi commodo, ipsum sed pharetra gravida, orci magna rhoncus neque, id pulvinar odio lorem non turpis. Nullam sit amet enim. Suspendisse id velit vitae ligula volutpat condimentum. Aliquam erat volutpat. Sed quis velit. Nulla facilisi. Nulla libero. Vivamus pharetra posuere sapien. Nam consectetur. Sed aliquam, nunc eget euismod ullamcorper, lectus nunc ullamcorper orci, fermentum bibendum enim nibh eget ipsum. Donec porttitor ligula eu dolor. Maecenas vitae nulla consequat libero cursus venenatis. Nam magna enim, accumsan eu, blandit sed, blandit a, eros.

It ain't what you don't know that gets you into trouble.

It's what you know for sure that just ain't so.

Commonly attributed to Mark Twain

1

Introduction

2

Background

This chapter introduces the concepts and tools used throughout this thesis. We begin with large language models and the hallucination problem they exhibit, then introduce the embedding spaces and geometric concepts that underpin our analysis. We conclude with the LLM-as-judge evaluation paradigm and the parameter-efficient fine-tuning methods used in our intervention experiments.

2.1 LARGE LANGUAGE MODELS

A large language model (LLM) is a neural network trained to predict the next token in a sequence, given all preceding tokens. The dominant architecture is the transformer [118], whose self-attention mechanism allows each token’s representation to be computed as a function of all prior tokens in the sequence. Through this mechanism, the model builds contextual representations of its input and produces a probability distribution over the vocabulary at each position. The training objective—

next-token prediction on large corpora of internet text—is a general-purpose objective that produces models capable of answering questions, summarizing documents, writing code, and generating coherent text across domains [74]. Crucially, this objective optimizes for *statistical plausibility* rather than factual accuracy directly: the model learns which tokens typically follow which contexts, and while factual statements are often the most plausible continuations, the objective provides no explicit mechanism for distinguishing true statements from confident-sounding fabrications.

Model capabilities improve predictably with parameter count, training data, and compute [97]. Some capabilities appear to emerge discontinuously at scale, including in-context learning (adapting to new tasks from a few examples in the prompt) and instruction following [122]. As a result, modern LLMs produce detailed, authoritative-sounding responses to nearly any question. The practical failure mode is not incoherence but *plausible fabrication*: the model generates text that reads as authoritative, yet is factually wrong.

INFERENCE. At inference time, the model generates tokens autoregressively. Each token is sampled from the predicted distribution, then appended to the context for the next prediction. A *temperature* parameter controls the sharpness of this distribution: temperature 0 selects the most probable token deterministically, while higher temperatures produce more diverse (and less predictable) outputs.

INSTRUCTION TUNING AND SYSTEM PROMPTS. Pre-trained LLMs generate plausible continuations but do not reliably follow user instructions. *Instruction tuning* [111] addresses this by fine-tuning the model on (instruction, response) pairs, teaching it to interpret and follow natural-language directives. Instruction-tuned models accept a *system prompt*: a prefix that conditions the model’s behavior for an entire conversation (e.g., “Answer concisely” or “If you are unsure, say so”). This system-prompt mechanism is the basis for the prompting interventions we study in Chapter 5.

OPEN-WEIGHT VS. CLOSED MODELS. Closed-source models (e.g., GPT-5.1, Claude Opus 4.5) are accessible only through APIs, with no access to weights or internal representations. Open-weight

models (e.g., Mixtral 8x7B, Llama 4 Maverick) provide full weight access, thus enabling fine-tuning and, in principle, extraction of internal representations. This distinction shapes our experimental design: our cross-model analysis (Chapter ??) includes both types, but the intervention experiments (Chapters 5 and 6) require open-weight models for LoRA fine-tuning (Section 2.5.2).

MIXTURE-OF-EXPERTS ARCHITECTURES. Both of our intervention target models use Mixture-of-Experts (MoE) routing, in which only a subset of parameters is active for each token. Mixtral 8x7B has 46.7 billion total parameters with approximately 12.9 billion active per token [93]; Llama 4 Maverick routes among 128 experts with approximately 17 billion active parameters per token. MoE architectures achieve strong performance at lower inference cost than equivalently-sized dense models, and their use of distinct experts for different inputs provides architectural diversity between our two target models (Chapter 4).

2.2 HALLUCINATION IN LARGE LANGUAGE MODELS

Hallucination in the LLM context refers to model-generated content that is fluent and confident but factually incorrect, fabricated, or unverifiable. The term is borrowed from cognitive science, where humans *confabulate*: producing plausible but false memories when probed at the boundaries of their knowledge. Unlike simple errors such as typos and misparsed inputs, LLM hallucinations are plausible-sounding statements generated with apparent confidence.

2.2.1 TAXONOMY

The standard taxonomy distinguishes *intrinsic* hallucinations, where the output contradicts a provided source document, from *extrinsic* hallucinations, where the output introduces information that cannot be verified from the source [92, 108]. This distinction was developed for tasks with explicit reference material (summarization, translation). In open-ended question answering—the setting studied in this thesis—there is no source document to contradict. The relevant question is not

how the output relates to a reference, but *why the model fabricated*: did it invent a nonexistent entity? Confuse a real entity with a similar one? Answer a question that has no definitive answer? This motivates our benchmark design (Chapter 4), which organizes prompts by the type of knowledge failure rather than by the source relationship.

2.2.2 CAUSES

Several factors contribute to hallucination in LLMs:

- Training objective mismatch. The next-token prediction objective assigns high probability to tokens that are *typical continuations* of the context. For well-documented entities, the correct fact dominates the training distribution and is reliably learned. But for rare or absent entities, the objective provides no mechanism to distinguish a correct claim from a plausible-sounding fabrication: both receive similar probability if they fit the syntactic and semantic context.
- Soft knowledge boundaries. LLMs do not have a hard boundary between “known” and “unknown.” Confidence degrades gradually with entity rarity: models reliably recall facts about common entities but produce increasingly unreliable outputs for entities in the long tail of the training distribution [115, 96]. This gray zone—where the model has partial, noisy, or absent training signal—is where hallucination concentrates.
- No explicit retrieval. Unlike a search engine, a standard (non-retrieval-augmented) LLM must reconstruct information from compressed parametric memory. When the memory signal for a query is weak, the model interpolates, generating text that is statistically plausible given the context, but not grounded in any specific source.
- Overconfidence. LLMs trained to be helpful tend to produce complete, confident-sounding answers even when their parametric knowledge is insufficient. Models can, to some degree, distinguish what they know from what they do not—Kadavath et al. [95] show that LLM

confidence correlates with accuracy—but they rarely express uncertainty spontaneously without explicit prompting [128].

2.2.3 CONSEQUENCES AND DETECTION CHALLENGES

Hallucinated content poses concrete risks in deployment: fabricated medical advice, invented legal citations, false historical claims. The core danger is that hallucinations are *plausible*—a user who cannot independently verify the claim may accept it as true, especially as overreliance on AI-generated text is a well-documented tendency in human-AI interaction [?]. As LLMs are deployed in search, customer service, education, and decision support, the surface area for hallucination-caused harm grows, and automated detection becomes necessary because human review does not scale.

Detection is difficult precisely because hallucinations are fluent and locally coherent. Surface-level features—grammar, length, apparent confidence—do not reliably distinguish hallucinated from correct responses. Two broad approaches exist: *post-generation detection*, which analyzes the model’s output or internal states after a response is produced [106, 83], and *pre-generation risk estimation*, which predicts *before generation* whether a given input is likely to elicit hallucination. This thesis explores the second approach through embedding geometry, while also investigating post-hoc mitigation via prompting and fine-tuning. Formal impossibility results suggest that hallucination cannot be fully eliminated from any LLM [127], motivating the focus on detection, prediction, and reduction rather than elimination.

2.3 TEXT EMBEDDINGS AND EMBEDDING GEOMETRY

The geometric analysis in this thesis operates on *text embeddings*—fixed-dimensional vector representations of natural language. This section introduces the embedding space, the distance metric we use, and the geometric concepts that our features (defined precisely in Chapter 4) formalize.

2.3.1 TEXT EMBEDDINGS

A text embedding model maps a variable-length token sequence to a fixed-dimensional real-valued vector $\mathbf{x} \in \mathbb{R}^d$. The mapping is learned via contrastive training: the model maximizes similarity between embeddings of semantically related texts and minimizes similarity between unrelated texts [113]. The result is a continuous vector space in which geometric proximity reflects semantic similarity: questions about similar topics are mapped to nearby points, while unrelated questions are mapped to distant ones [72].

2.3.2 COSINE DISTANCE

All distance computations in this thesis use *cosine distance*. For two vectors $\mathbf{x}, \mathbf{y} \in \mathbb{R}^d$, the cosine similarity is

$$\text{sim}(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}, \quad (2.1)$$

and the cosine distance is $d_{\text{cos}}(\mathbf{x}, \mathbf{y}) = 1 - \text{sim}(\mathbf{x}, \mathbf{y})$, ranging from 0 (identical direction) to 2 (opposite direction). Cosine distance is invariant to vector magnitude and depends only on the angle between vectors, making it appropriate for embeddings where direction carries more semantic information than magnitude. When vectors are L2-normalized (as in our setting), cosine distance and squared Euclidean distance are related by $\|\mathbf{x} - \mathbf{y}\|^2 = 2(1 - \text{sim}(\mathbf{x}, \mathbf{y}))$, so the choice of metric is equivalent.

2.3.3 HIGH-DIMENSIONAL GEOMETRY

Our embeddings reside in $\mathbb{R}^d$ for $d = 3,072$ (the dimensionality of the embedding model selected in Chapter 4). High-dimensional spaces exhibit counterintuitive geometric properties that affect how we interpret distance-based features. A well-known example is *distance concentration*: pairwise distances between random points in high dimensions converge to a narrow range, making absolute distances less informative [67]. Within structured datasets, however, relative distances—which points are closer or farther than others—remain meaningful. All of our geometric features (Chapter 4) are

defined relative to the corpus distribution rather than in absolute terms, making them robust to concentration effects. When embeddings are L2-normalized, all points lie on the surface of the unit hypersphere in $\mathbb{R}^{3072}$, and cosine distance is the natural metric on this surface.

2.3.4 THE MANIFOLD HYPOTHESIS

High-dimensional data often lies on or near a lower-dimensional *manifold*—a smooth subspace embedded in the ambient space [84]. Text embeddings are not uniformly distributed across $\mathbb{R}^{3072}$; they concentrate on a structured subspace determined by the distribution of natural language. This structure creates meaningful local geometry: some regions of the manifold are densely populated (many semantically similar texts), while others are sparse; the manifold may curve differently in different regions; and certain directions of variation are more prominent than others.

The central hypothesis of this thesis, stated informally, is that geometric properties of a prompt’s position on this manifold—its local density, the structure of its neighborhood, its distance from typical regions—carry information about whether a language model will hallucinate in response. This is a hypothesis, not an assumption; Chapter ?? tests it empirically.

2.3.5 GEOMETRIC CONCEPTS

The geometric features defined in Chapter 4 each formalize a different aspect of manifold structure. Before those precise definitions, we introduce the underlying concepts.

***k*-NEAREST NEIGHBORS.** For a point x_i in a dataset of N points, its k -nearest neighbors (k -NN) are the k points with smallest distance to x_i under the chosen metric. The k -NN set is the building block for most of our geometric features: it defines the “local neighborhood” within which we measure density, dimensionality, and curvature. The choice of k controls the locality of the analysis: small k captures fine-grained structure; large k smooths over broader neighborhoods.

PRINCIPAL COMPONENT ANALYSIS. Principal Component Analysis (PCA) finds the orthogonal directions of maximum variance in a dataset—the eigenvectors of the covariance matrix, ordered by

eigenvalue magnitude. The *explained variance ratio* of the top p components,

$$\text{EVR}(p) = \frac{\sum_{j=1}^p \sigma_j^2}{\sum_{\ell} \sigma_{\ell}^2}, \quad (2.2)$$

where σ_j^2 are the eigenvalues, measures how well a p -dimensional linear subspace approximates the data. PCA appears in two roles in our work: measuring how well a flat subspace approximates a local neighborhood (curvature), and constructing a novel “semantic opposite” of each point (oppositeness). Both are defined precisely in Chapter 4.

INTRINSIC DIMENSIONALITY. The *intrinsic dimension* of a dataset at a point estimates the number of effective directions of variation in the local neighborhood. A manifold embedded in 3,072-dimensional space may have far fewer effective local dimensions. Standard estimators include the TwoNN method [82], which uses only the ratio of distances to the two nearest neighbors, and the maximum likelihood estimator of Levina & Bickel [98], which averages log-distance ratios over a larger neighborhood for a smoother estimate.

LOCAL DENSITY. The *density* at a point measures how many similar points surround it—equivalently, the inverse of the average distance to nearby neighbors. High density indicates a well-represented region with many semantically similar texts; low density indicates an isolated point with few nearby neighbors, potentially representing out-of-distribution or under-represented content. Density-based measures are standard in anomaly detection [73].

CURVATURE. The *curvature* of the manifold at a point reflects how well a flat (linear) subspace approximates the local neighborhood. If the neighborhood lies on a flat subspace, PCA captures most of the variance and the residual is small (low curvature). If the neighborhood bends or twists, the linear approximation misses variance and the residual is large (high curvature).

CENTRALITY. The *centrality* of a point, as used in this thesis, measures its distance from the center of the distribution (e.g., the corpus mean embedding); higher values indicate more *peripheral* points. Points far from the center may represent unusual, specialized, or boundary-region content.

OPPOSITENESS. In addition to these standard geometric concepts, Chapter 4 defines a novel *oppositeness* feature that measures how far a point’s “semantic opposite” (constructed by flipping principal component directions) lies from the nearest real data point. The construction is specific to our experimental design and is defined precisely in Section 4.4.3.

2.4 EVALUATING LANGUAGE MODEL OUTPUTS

Evaluating whether an LLM response is factually correct, partially correct, hallucinated, or a refusal requires semantic judgment that scales beyond traditional automated metrics and human annotation. This section introduces the LLM-as-judge paradigm that our evaluation pipeline builds on.

2.4.1 THE EVALUATION BOTTLENECK

A typical hallucination study may require evaluating tens of thousands of model responses across multiple models and experimental conditions. Human annotation at this scale is infeasible. Traditional automated metrics such as BLEU and ROUGE measure surface-level overlap with reference texts and cannot capture the semantic judgment needed for hallucination detection—whether a response fabricates an entity or invents a citation cannot be determined by n -gram matching.

2.4.2 LLM-AS-JUDGE

The LLM-as-judge approach uses a large language model to evaluate the outputs of another LLM. The judge receives the original question, the model’s response, and optionally a ground-truth reference, then produces a label and confidence score according to a rubric. Zheng et al. [130] demonstrated that GPT-4 judgments align with human preferences over 80% of the time on conversational

quality assessment, and Liu et al. [103] showed that rubric-based LLM evaluation outperforms traditional metrics on summarization quality. Two modes are common: *reference-free* evaluation, where the judge assesses quality without a gold standard, and *reference-based* evaluation, where the judge compares the response against a provided ground truth. Our pipeline uses reference-based evaluation, providing judges with the known correct answer to reduce subjectivity.

2.4.3 KNOWN BIASES

LLM judges exhibit several documented systematic biases. *Verbosity bias*: longer, more detailed responses tend to receive higher ratings even when the additional content is not more accurate [130]. *Position bias*: in pairwise comparisons, the response presented first or second may receive systematically different treatment [130]. *Self-preference bias*: a model may rate its own outputs more favorably than those of other models [121, 112]—relevant in settings where evaluated models also serve as judges. *Shared blind spots*: judges trained on similar data may share systematic errors, all failing to detect the same hallucination.

These biases motivate two design principles for our evaluation pipeline (Chapter 4): using a *multi-judge panel* rather than a single judge (reducing any individual judge’s influence), and selecting judges from *different model families* (reducing correlated errors). The cost of a panel scales linearly with the number of judges; an odd-numbered panel ($n = 3$) avoids ties in majority-vote aggregation while remaining affordable.

2.4.4 RUBRIC DESIGN

Hallucination evaluation requires distinguishing among multiple response types. A binary correct/incorrect rubric loses the distinction between a model that fabricates content (active harm) and one that refuses to answer (safe but unhelpful). This distinction matters for fine-tuning (Chapter 6): fabricated responses should be excluded from training data, while refusals and partial responses carry different training signal. Multi-category rubrics are standard in hallucination evaluation—TruthfulQA uses separate truthfulness and informativeness axes [102]; our four-point rubric is de-

finetuned in Chapter 4.

2.5 FINE-TUNING AND PARAMETER-EFFICIENT ADAPTATION

Chapter 6 presents experiments in which we fine-tune language models to reduce hallucination. This section introduces the relevant concepts.

2.5.1 SUPERVISED FINE-TUNING

After pre-training on next-token prediction, a model can be *fine-tuned* on a curated dataset of (input, desired output) pairs to adapt its behavior for a specific task. *Instruction tuning* [111] is a widely-used variant in which the fine-tuning data consists of instructions paired with high-quality responses. Fine-tuning trades generality for task-specific performance, with the risk of *overfitting*—memorizing specific training answers rather than learning generalizable behavior. Even modest datasets can be effective: Zhou et al. [131] demonstrate that fine-tuning on as few as 1,000 carefully curated examples can substantially improve instruction following, suggesting that data quality matters more than quantity.

2.5.2 LoRA: LOW-RANK ADAPTATION

Full fine-tuning updates all model parameters, which is prohibitively expensive for models with tens or hundreds of billions of parameters. *Low-Rank Adaptation* (LoRA) [90] instead freezes the pre-trained weights and adds small trainable matrices to the attention layers. For a weight matrix $W \in \mathbb{R}^{d \times d}$, LoRA parameterizes the update as

$$\Delta W = BA, \quad B \in \mathbb{R}^{d \times r}, \quad A \in \mathbb{R}^{r \times d}, \quad r \ll d, \quad (2.3)$$

so the effective weight at inference is $W + BA$. The rank r controls the capacity of the adaptation: typical values range from 8 to 64, reducing trainable parameters to roughly 0.1% of the full model. A scaling factor α controls the effective learning rate of the adapter (the update is scaled by

α/r). LoRA inherently limits the risk of catastrophic forgetting—overwriting the model’s general capabilities—because the vast majority of weights remain frozen.

QLoRA [79] extends LoRA by quantizing the frozen base model to 4-bit precision, further reducing memory requirements while maintaining fine-tuning quality. Our fine-tuning experiments (Chapter 6) use LoRA for Mixtral and QLoRA for Llama 4 Maverick, as determined by the fine-tuning API.

2.5.3 SELF-DISTILLATION

Our fine-tuning approach (Chapter 6) is a form of *self-distillation*: the model is fine-tuned on its own best outputs, selected across multiple prompting conditions. Rather than training on externally-generated targets, the model learns to produce its best observed behavior by default, without requiring a system prompt at inference time. Self-distillation is well-established as a technique for improving model performance: Hinton et al. [89] introduced knowledge distillation between models, and subsequent work showed that training on a model’s own outputs can improve performance through implicit regularization [86].

This chapter has introduced the conceptual foundations for the remainder of this thesis: the structure and generation process of large language models, the hallucination problem and its causes, the embedding spaces and geometric concepts that underpin our analysis, the LLM-as-judge evaluation paradigm, and parameter-efficient fine-tuning. Chapter 3 surveys the existing literature on hallucination detection, mitigation, and geometric analysis of neural representations, identifying the specific gaps that this thesis addresses. Chapter 4 then describes our experimental infrastructure.

Geometry is not true, it is advantageous.

Henri Poincaré

3

Literature Review

This chapter surveys prior work across four areas that converge in this thesis: hallucination taxonomy, benchmarking, and evaluation methodology; representation-based approaches to truthfulness; entity knowledge and its limits; and intervention via prompting and fine-tuning. Each section identifies a specific gap that motivates the research questions addressed in later chapters. A concluding synthesis distills these gaps into the three questions that structure the remainder of the thesis.

3.1 HALLUCINATION: TAXONOMY AND THEORETICAL LIMITS

Maynez et al. [108] established a foundational distinction between *intrinsic* hallucinations, which contradict a source text, and *extrinsic* hallucinations, which introduce unverifiable information. Subsequent surveys generalized this taxonomy across natural language generation (NLG) tasks. Ji et al. [92] mapped causes, definitions, and mitigations across modalities; Huang et al. [91] introduced a

distinct axis of *factuality* hallucination (claims contradicting world knowledge) versus *faithfulness* hallucination (claims contradicting a provided context); and Tonmoy et al. [117] compiled the most detailed mitigation comparison to date, covering 32 methods. This thesis focuses on factuality hallucination: cases where a model generates claims about entities or facts that are verifiably false.

A key theoretical result motivates the shift from elimination to prediction. Xu et al. [127] proved that hallucination cannot be fully eliminated: any computable LLM will hallucinate on some inputs. If total prevention is impossible, the research agenda must turn to *predicting where* hallucinations will occur and *mitigating the ones that do*. This is the central premise of our geometric prediction approach.

Yet existing surveys treat hallucination as a behavioral property of model outputs. None ask whether hallucination has a geometric *signature* in the representation space of the input—a signal that could be detected before generation begins.

3.2 BENCHMARKS FOR FACTUALITY EVALUATION

Factuality benchmarks have evolved through several generations, each probing a deeper aspect of the problem. Early efforts focused on *whether* models hallucinate: TRUTHFULQA [102] tested whether models reproduce common human falsehoods, producing the surprising finding that larger models are not necessarily more truthful, and HELM [101] incorporated factuality into holistic evaluation across capabilities and risks. Subsequent benchmarks shifted toward *how much* and *at what granularity*: HALUEVAL [99] scaled to 35,000 samples across QA, dialogue, and summarization, while FACTSCORE [109] decomposed biography generation into atomic claims, revealing that only 58% of ChatGPT’s biographical statements are supported by knowledge sources. More recent work such as SIMPLEQA [123] tests short-form factual recall with unambiguously verifiable answers, tightening the evaluation loop further.

A parallel line of benchmarks asks not just how often models fail, but *on what kinds of entities*. POPQA [105] organized 14,000 questions by Wikipedia entity popularity and established that LLM accuracy correlates directly with training data frequency—a result that reframes hallucination as

partly a data coverage problem rather than a purely architectural one. Our benchmark complements this popularity-based stratification by introducing a distinct axis: *existence status*—real, nonexistent, borderline—testing failure modes at the boundary of model knowledge rather than across the popularity spectrum (Chapter 4).

Despite this progression, existing benchmarks quantify *how often* models hallucinate and *on what topics*, but none are designed to enable analysis of *why* certain prompts are harder in geometric terms. No benchmark stratifies by representational properties of entities in embedding space.

3.3 AUTOMATED EVALUATION AND THE LLM-AS-JUDGE PARADIGM

Human evaluation remains the gold standard but does not scale. Zheng et al. [130] demonstrated with MT-Bench and Chatbot Arena that LLM judges achieve strong human–model agreement, especially when diverse architectures are used in consensus. G-Eval [103] showed that GPT-4 outperforms traditional NLG metrics. The Panel of LLM evaluators (PoLL) approach [119] further showed that panels of smaller, diverse-family models outperform a single large judge while costing less and exhibiting lower intra-model bias—directly motivating our three-family consensus design.

Known biases in LLM judges have been extensively characterized. Self-preference bias is real and causal: Panickssery et al. [112] demonstrated that LLMs can identify their own outputs and rate them higher, with a linear relationship between self-recognition ability and bias strength. Wataoka et al. [121] confirmed this across model families. These findings motivate using judges from *different* model families than the models being evaluated, and aggregating via majority vote to dilute individual biases.

Multi-judge consensus has been validated for preference ranking and general-purpose NLG evaluation, but its application to *hallucination-specific* multi-label classification (correct, partial, hallucinated, refused) on entity-level prompts has received little study.

3.4 DETECTING HALLUCINATION FROM INTERNAL REPRESENTATIONS

A rapidly growing body of work uses internal model representations to detect hallucination. This section is the most directly relevant to our thesis, and we review it in detail.

3.4.1 PROBING FOR TRUTHFULNESS

Burns et al. [76] showed with Contrast-Consistent Search (CCS) that truth directions exist in LLM hidden states without any supervision—models often “know” the correct answer internally even when generating incorrectly. Azaria & Mitchell [70] trained classifiers on hidden-layer activations, detecting truthfulness with 71–83% accuracy. Marks & Tegmark [107] found that LLMs linearly represent truth and falsehood: linear probes trained on one factual dataset generalize to others, and causal interventions along the truth direction flip model judgments. This “Geometry of Truth” work is the most direct conceptual antecedent to our thesis.

However, subsequent work has complicated the picture. Orgad et al. [110] showed that truthfulness encoding is *not universal*—it generalizes only within tasks requiring similar cognitive skills, challenging the idea of a single truth direction. This is an important qualifier for our own cross-category results: geometric features that predict hallucination within one entity category may not transfer to categories requiring different skills.

3.4.2 GEOMETRIC AND SPECTRAL METHODS

Where probing methods train classifiers on raw activations, a parallel line of work asks whether the *geometry* of representation space itself carries truthfulness information—shifting from “can we learn to detect lies?” to “does the shape of the manifold encode them?”

Singhal et al. [114] provided early evidence that it does: they showed that Local Intrinsic Dimension (LID) of model activations predicts hallucination, with truthful outputs occupying more structured, lower-dimensional manifold neighborhoods. Subsequent methods exploited this geometric structure in different ways. HaloScope [81] identifies a hallucination-associated *subspace* via SVD

factorization, matching supervised detection accuracy without labels—suggesting the truthfulness signal is concentrated in a low-rank component of the activation space. INSIDE [77] takes a spectral approach, computing eigenvalues of response embedding covariance matrices and outperforming logit-based baselines. Most strikingly, Burger et al. [75] found that a single two-dimensional subspace separates true from false across multiple LLM families (Gemma, LLaMA, Mistral), suggesting that geometric structure for truthfulness may be shared across architectures rather than being model-specific.

3.4.3 UNCERTAINTY-BASED AND INFERENCE-TIME METHODS

The methods above detect hallucination from static representations; a separate class of approaches instead exploits the *dynamics* of the generation process itself. Sampling-based methods probe output consistency: Semantic Entropy [83] clusters semantically equivalent sampled answers before measuring entropy, achieving state-of-the-art detection across five QA datasets, while SelfCheck-GPT [106] uses a simpler consistency check requiring no external knowledge. Both rely on the insight that a model confident in a fact will produce consistent answers, while hallucinated claims vary across samples.

A more interventionist line of work goes beyond detection to *causally steering* model outputs. ITI [100] identifies truthful directions in attention head activations and shifts representations along them, roughly doubling truthfulness scores on TruthfulQA. Representation Engineering [132] generalizes this to population-level steering of truth-related features. DoLA [78] takes a different approach entirely, contrasting logit distributions across early and late transformer layers to amplify factual knowledge that is localized in specific layers. Collectively, these methods establish that truthfulness information not only exists in model representations but can be causally manipulated—a finding that motivates asking, as we do, whether similar structure exists in the *input* embedding space.

All existing work detects hallucination from *internal* model representations *during or after* generation—requiring white-box access to activations, hidden states, or sampling distributions. Our

approach differs in two ways: (1) we predict hallucination from the *embedding geometry of the input prompt*, before generation, using an external embedding model with no white-box access to the target model; and (2) we predict not only hallucination *occurrence* but hallucination *difficulty*—which hallucinations resist mitigation and why. No prior work addresses fixability prediction.

3.5 ENTITY KNOWLEDGE AND KNOWLEDGE BOUNDARIES

Kandpal et al. [96] established the foundational result: LLM accuracy on factual questions correlates with the number of documents mentioning the relevant entity in training data. Scaling model size improves recall of popular entities but fails on the long tail. Mallen et al. [105] confirmed this with PopQA and showed that retrieval augmentation helps rare entities but can hurt popular ones. Sun et al. [115] demonstrated with Head-to-Tail that hallucination rate is inversely correlated with entity popularity across knowledge graph triples.

At the mechanistic level, Ferrando et al. [85] applied sparse autoencoders to LLM residual streams and discovered *entity recognition features*—internal directions that detect whether a model can recall facts about a given entity. These features generalize across entity types and are causally relevant: activating them can induce hallucination about known entities or refusal about unknown ones. Kadavath et al. [95] showed models can predict $P(\text{IK})$ —the probability of knowing an answer—establishing that LLMs maintain internal representations of their own knowledge boundaries. Yin et al. [128] confirmed that models possess self-knowledge capacity but with a significant gap versus human proficiency.

Entity popularity is a well-established predictor of hallucination, but its connection to *embedding geometry* has not been tested. Our density metric may capture a related signal: entities in sparse embedding neighborhoods may correspond to entities underrepresented in training data, where models default to fabrication over appropriate refusal. Our benchmark is specifically designed with entity-existence categories—nonexistent, plausible-but-fake, obscure-but-real—that test this geometric–knowledge boundary hypothesis.

3.6 PROMPT ENGINEERING FOR FACTUALITY

Prompting-based hallucination mitigation spans reasoning augmentation and system-level instruction design. The dominant paradigm has been multi-pass reasoning. Chain-of-thought prompting [124] dramatically improved LLM reasoning, though its effect on factuality specifically is mixed. Subsequent methods built on this foundation: Self-Consistency [120] samples multiple reasoning paths and takes a majority vote, Chain-of-Verification [80] adds explicit self-verification stages that reduce factual hallucination by 50–70%, and Self-Refine [104] iterates on outputs using self-generated feedback. These approaches share a common tradeoff: they improve reliability at the cost of multiple generation passes per query, making them expensive at inference time.

At the system-prompt level, fewer controlled studies exist. SynTra [94] is the closest prior work to ours: it optimizes continuous prefix embeddings for system messages via synthetic task transfer, reducing hallucination by 29%. However, SynTra requires gradient-based optimization, while our discrete, interpretable prefixes need no training. Arora et al. [68] showed that explicit, structured prompting rules materially improve truthfulness over vague instructions. The abstention literature provides additional context: Wen et al. [125] survey abstention approaches comprehensively, noting that fine-tuning can degrade a model’s willingness to abstain, motivating prompt-based abstention strategies as a complementary approach. R-Tuning [129], which teaches refusal as a transferable meta-skill via fine-tuning, demonstrates that the abstention capability our prompts attempt to activate does generalize beyond training distributions.

No prior work systematically compares multiple discrete system-prompt-level interventions on the same entity-level hallucination benchmark under controlled conditions. SynTra uses continuous optimization; CoVe requires multi-stage pipelines; most studies test a single prompting strategy. Furthermore, no work connects prompt effectiveness to geometric properties of the prompts being modified—our bridge analysis (Chapter 6) tests exactly this.

3.7 FINE-TUNING FOR TRUTHFULNESS

3.7.1 ALIGNMENT AND ITS COSTS

InstructGPT [111] showed that RLHF improves instruction-following but does not fully address factuality. Constitutional AI [71] uses self-critique and revision guided by explicit principles. A growing body of evidence suggests alignment can actively *hurt* factuality: standard supervised fine-tuning on human-written answers introduces knowledge the model does not have, increasing hallucination [87]. This “fine-tuning paradox” implies that SFT should teach behavioral patterns, not inject new factual knowledge—a principle central to our pipeline design.

3.7.2 FACTUALITY-TARGETED FINE-TUNING

R-Tuning [129] is the closest comparator to our fine-tuning approach. Both teach models to handle knowledge boundaries, but through different signals: R-Tuning uses binary train-time probing (“can the model answer this?”), while our pipeline selects training data by testing multiple prompt interventions and keeping the best response per prompt. Our geometric analysis then explains *why* certain prompts resist correction—a diagnostic role, not an operational one. Other approaches differ in mechanism: FactTune [116] uses DPO with automatically generated preference pairs, achieving a 58% reduction in factual errors on biography generation, while Self-RAG [69] combines fine-tuning with learned retrieval to ground generation in evidence.

3.7.3 DATA SELECTION

Evidence increasingly shows that training data quality dominates quantity. LIMA [131] demonstrated that 1,000 carefully curated examples match much larger datasets for alignment. LESS [126] showed that influence-function-selected 5% of data often outperforms the full dataset, and that selected data transfers across model families. These results motivate our best-per-prompt selection strategy, which curates training data by choosing, for each prompt, the response generated under the most effective intervention.

Existing fine-tuning approaches select training data by model performance (R-Tuning), preference pairs (DPO family), or influence functions (LESS). None analyze the geometric properties of the resulting training distribution—embedding density, centrality, oppositeness—to understand *why* certain prompts are harder to correct. Our bridge analysis fills this diagnostic gap, showing that geometric features predict which hallucinations resist intervention (Chapter 6). Furthermore, no prior work quantifies the precision–recall tradeoff between hallucination reduction and false refusal of obscure-but-real entities—a tradeoff we document in that chapter.

3.8 SYNTHESIS AND RESEARCH QUESTIONS

The literature reviewed above reveals three convergent gaps.

A geometric gap. Section 3.4 showed that geometric properties of representations encode truthfulness, and Section 3.5 showed that entity frequency predicts hallucination. But no work has tested whether geometric features of *prompt embeddings*—density, oppositeness, centrality, curvature—predict hallucination before generation and across models.

An intervention selection gap. Section 3.6 showed that prompting reduces hallucination, and Section 3.7 showed that fine-tuning can teach abstention. But which prompts benefit from which intervention—and can geometric features predict this?

A distillation gap. No work has tested whether the behavioral effects of prompt-based interventions can be consolidated into model weights via fine-tuning, with geometric analysis explaining where the resulting model succeeds and fails.

These gaps motivate three research questions:

1. Geometric prediction (Chapter ??): Can geometric properties of prompt embeddings predict hallucination, and does this signal survive after controlling for category structure?
2. Prompt intervention (Chapter 5): Can discrete system-prompt prefixes reduce hallucination in a controlled comparison, and do different prefixes work for different entity categories?

3. Geometry-guided fine-tuning (Chapter 6): Can geometric features predict which hallucinations are fixable by prompting, and can the best prompt behavior be distilled into model weights?

Together, these questions form a progressive narrative: geometry as diagnostic, prompts as treatment, fine-tuning as consolidation—with geometry predicting where the consolidation has side effects. The literature reviewed in this chapter shows each component exists in isolation; the contribution of this thesis is connecting them into an integrated, empirically validated pipeline.

4

Experimental Setup

This chapter describes the experimental infrastructure shared across all experiments in this thesis: the hallucination benchmark, the language models under study, the consensus judging pipeline that converts model responses into hallucination labels, and the embedding approach that enables geometric analysis. Experiment-specific methods—geometric feature extraction, prompt prefix design, and fine-tuning procedures—are introduced in their respective chapters. Figure 4.1 provides a high-level overview of the shared infrastructure.




Figure 4.1: Overview of the shared experimental infrastructure. Prompts from the benchmark are sent to language models for generation, responses are evaluated by a three-judge consensus panel, and prompt embeddings are analyzed for geometric features. Each component corresponds to a section in this chapter.

4.1 BENCHMARK CONSTRUCTION

Evaluating hallucination requires a benchmark that systematically probes the boundary between what a model knows, and what it does not. Existing hallucination benchmarks such as TruthfulQA [102] and SimpleQA [123] focus primarily on factual recall, testing whether models produce correct answers to verifiable questions. While valuable, these benchmarks do not control for the *type* of knowledge failure: for instance, a model that fabricates a nonexistent scientist’s biography fails differently from one that confidently answers an inherently ambiguous question, and both of these fail differently from one that hallucinates specific details about a real but obscure historical figure.

Our benchmark is designed to probe these distinctions. We construct prompts across seven categories that span the space of possible knowledge failures, from straightforward factual questions (where hallucination reflects a clear error) to borderline cases (where even the definition of “correct” is contested). Each category targets a distinct failure mode, and the category structure itself becomes an object of analysis: in Chapter ??, we investigate whether geometric properties of prompts in embedding space vary systematically across categories, and whether this variation relates to hallucination rates.

The benchmark exists in two versions. The initial benchmark comprises 449 prompts and serves as a held-out test set throughout this thesis. The expanded benchmark adds 2,430 new prompts with zero overlap, designed for fine-tuning training data with greater template and entity diversity. Both versions share the same category structure, generation methodology, and evaluation pipeline.

4.1.1 CATEGORY DESIGN

We define seven prompt categories, each targeting a distinct mechanism by which a language model might hallucinate. Prior hallucination taxonomies distinguish between *intrinsic* hallucinations, where the output contradicts the source material, and *extrinsic* hallucinations, where the output introduces information that cannot be verified from the source [92, 108]. However, this distinction was designed for tasks with a reference document (e.g., summarization). In open-ended generation,

there is no source to contradict—the relevant question is not how the output relates to a nonexistent reference, but *why* the model hallucinated. Our categories are organized along two axes that address this: the *type of knowledge* being tested (entity existence, factual recall, logical reasoning, subjective judgment) and the *expected difficulty* (from unambiguous to deliberately challenging).

FACTUAL. Prompts with unambiguous, objectively verifiable answers grounded in established knowledge. Examples include “What is the chemical formula for sodium chloride?” and “Who wrote the book *1984*?” These serve as a baseline control, analogous to the factual recall questions in TruthfulQA [102] and HaluEval [99]. Hallucination on factual prompts reflects a clear failure of knowledge recall.

NONEXISTENT. Prompts referencing entities, works, or concepts that do not exist. Examples include “Summarize the book *The Obsidian Codex* by Dr. Marcus Valmont” and “Explain the Bertrand-Goldstein theorem in mathematics.” The correct response is to identify the entity as nonexistent; any detailed description of the entity constitutes hallucination. This category operationalizes the question of whether models “know what they don’t know” [95, 128]. It is the largest category in the benchmark (Table 4.1) because it targets the failure mode most central to hallucination: unlike factual prompts, where correct information may be available in the model’s training data, nonexistent prompts offer no correct content to recall—any detailed response necessarily constitutes fabrication.

IMPOSSIBLE. Prompts requesting tasks that cannot be completed by any system: proving unsolved conjectures, predicting future events with certainty, reading the user’s thoughts, or listing infinite sets. Examples include “Provide a complete proof of the Birch and Swinnerton-Dyer conjecture” and “Tell me exactly what I’m thinking right now.” Unlike nonexistent prompts, which test entity knowledge, impossible prompts test whether models recognize the *logical or physical* limits of answerable questions.

AMBIGUOUS. Prompts with no single correct answer, where the response legitimately depends on subjective criteria, values, or framing. Examples include “What makes a good leader?” and “Is it foolish to lie to protect someone?” The correct behavior is to acknowledge the ambiguity rather than present a subjective opinion as fact. Fabricating verifiable claims (e.g., citing a nonexistent study) within an otherwise subjective response still constitutes hallucination. The precise boundary between acceptable subjectivity and hallucination is operationalized through the consensus judging pipeline described in Section 4.3.

BORDERLINE: PLAUSIBLE FAKE. Entities that sound real but do not exist—the adversarial counterpart to nonexistent prompts, probing the region just beyond the tail of Sun et al. [115]’s entity popularity spectrum. Where nonexistent entities are often transparently fabricated (e.g., the country “Quantia,” the disease “Chromatic Resonance Fatigue”), plausible fakes are designed to be difficult to distinguish from real entities: “Dr. Helena Vasquez-Okonkwo,” “*The Lintel Above the Dovecote*,” “Penderwych Narrows.” We expect this to be the hardest category in our benchmark: unlike transparently fake entities, plausible fakes cannot be rejected on surface features alone, forcing the model to rely on genuine knowledge of whether the entity exists.

BORDERLINE: OBSCURE REAL. Real entities that are sufficiently obscure that a model might incorrectly deny their existence—the “tail” entities in the terminology of Sun et al. [115], where model knowledge degrades sharply. Examples include Rosalind Franklin (who is often underattributed), Tristan da Cunha (the world’s most remote inhabited island), and the Cadaver Synod (a medieval papal trial of a corpse). These prompts test the opposite failure mode from plausible fakes: *false negatives* on entity existence, where a model refuses to answer about something that genuinely exists.

BORDERLINE: EDGE FACTUAL. Factual questions designed to elicit confidently wrong answers—commonly confused facts, counter-intuitive truths, and widespread misconceptions. Examples include “What is the largest desert in the world?” (Antarctica, not the Sahara), “What is the capital of Turkey?” (Ankara, not Istanbul), and “What is the probability of two people in a group of 23

sharing a birthday?” (approximately 50.7%). Unlike standard factual prompts, edge factual prompts target knowledge boundaries where models are likely to reproduce common human errors—what Lin et al. [102] term “imitative falsehoods,” where models learn to repeat popular misconceptions from training data. Each prompt is individually curated with a verified ground truth rather than generated from templates.

Table 4.1 summarizes the seven categories with their sizes and expected correct behavior.

Category	Test	Train	Total	Expected Behavior
Factual	98	500	598	Answer correctly
Nonexistent	120	600	720	Identify as nonexistent
Impossible	30	200	230	Acknowledge impossibility
Ambiguous	120	600	720	Acknowledge subjectivity
Borderline: Plausible Fake	31	200	231	Identify as nonexistent
Borderline: Obscure Real	30	200	230	Answer or acknowledge uncertainty
Borderline: Edge Factual	20	130	150	Answer correctly (not the common error)
Total	449	2,430	2,879	

Table 4.1: Benchmark categories with sizes across the held-out test set (449 prompts) and expanded training set (2,430 prompts). The seven categories span the space of knowledge failures from straightforward factual recall to adversarial borderline cases. The two sets have zero prompt overlap.

WHY SEVEN CATEGORIES. The taxonomy was defined a priori based on the axes of knowledge failure that were identified in the literature, then held fixed throughout all experiments—no categories were added, removed, or redefined after data collection began. The four main categories (factual, nonexistent, impossible, ambiguous) cover the primary axes of knowledge failure identified in prior surveys [92]: incorrect recall, entity fabrication, reasoning limits, and subjectivity. The three borderline categories were added to probe the decision boundary between “knows” and “does not know”—the region where models are most uncertain and where we hypothesize geometric features of embedding space to be most informative. A coarser taxonomy (e.g., “factual” versus “non-factual”) would collapse distinctions that matter: a model that hallucinates on plausible fakes but not on transparent fakes has a qualitatively different failure mode than one that hallucinates

uniformly. A finer taxonomy (e.g., separate categories for each entity subtype) would fragment the data without adding discriminative power, as we investigate in Chapter ??, where we test whether geometric features vary primarily across the seven categories rather than within subtypes.

4.1.2 TEMPLATE DESIGN

Prompts are generated from *templates*: structural patterns with placeholder variables that are filled from entity pools at generation time. For example, the template "What is the chemical formula for {compound}?" combined with the entity pool `compound = [water, sodium chloride, . . .]` produces concrete prompts like "What is the chemical formula for water?"

Templates serve two purposes. First, they *control for syntactic variation*: by holding sentence structure constant and varying only the entity, we ensure that differences in model behavior across prompts within a category reflect differences in entity knowledge rather than differences in how the question is phrased. Second, they enable *scalable generation*: from 55–66 templates per main category and 15–44 entities per pool, we can generate hundreds of unique prompts with combinatorial diversity while maintaining structural consistency.

TEMPLATE DIVERSITY. Each main category (factual, nonexistent, impossible, ambiguous) has 55–66 templates covering diverse question formats: direct factual questions ("What is...?"), existence probes ("Who is...?"), comparison questions ("How does X compare to Y?"), true/false framings ("True or false: ..."), and open-ended requests ("Describe...", "Explain..."). Borderline categories have 20 templates per subtype (6 subtypes $\times$ 20 = 120 borderline templates total). Templates were designed to elicit diverse response patterns: some invite short factual answers, while others encourage extended generation where hallucination is more likely.

TEMPLATE COVERAGE. To ensure no single template dominates the benchmark, prompts are generated using *round-robin sampling*: templates are shuffled and cycled through sequentially, with reshuffling at each cycle boundary. This guarantees that each template is used approximately equally.

In the expanded benchmark, template coverage reaches 100% for factual, nonexistent, and ambiguous categories (all templates used at least once) and 69.7% for impossible. This is expected: 20 of 66 impossible templates are static single-use prompts (e.g., “Prove that $P = NP$ ”) that appeared verbatim in the initial benchmark; because the expansion enforces zero overlap with the held-out test set, these templates are excluded from the training set.

ENTITY FILLING. When a template contains multiple placeholders (e.g., “Who is {fake_name}, the inventor of {fake_invention}?”), each variable is filled independently from its respective entity pool. For paired variables that must be semantically coherent (e.g., {option1} and {option2} in comparison templates), we use structured entity pools where each entry contains related values that are resolved together.

PLACEHOLDER VALIDATION. Every generated prompt is checked for unfilled placeholders (patterns matching `[{variable}]`). Any prompt with residual placeholders is rejected and regenerated. The initial benchmark contained approximately 20 prompts with unfilled placeholders that were not caught; the expanded benchmark enforces this validation strictly, resulting in zero unfilled placeholders across all 2,430 prompts.

4.1.3 ENTITY POOLS AND VALIDATION

Entity pools are the knowledge substrate of the benchmark. Their quality directly determines whether hallucination labels are meaningful: a “nonexistent” entity that turns out to be real produces a false hallucination label, and a “factual” entity with an incorrect ground truth produces a false correctness label.

SOURCING. Entity pools are constructed differently by category:

- Factual: Entities drawn from established knowledge domains—countries, chemical compounds, literary works, historical events, scientific facts. Each entity has a verifiable ground

truth. Pools range from 15 to 44 entities (e.g., 44 countries, 15 chemical compounds, 15 canonical books).

- **Nonexistent:** Entities manually crafted to be recognizably fictional. Names follow conventions within their domain (e.g., “Bertrand-Goldstein theorem” for mathematics, “Zamboria” for countries) but do not correspond to real entities. 34 entity pools with 514 total entities.
- **Impossible:** Pools of unsolved conjectures, inherently unpredictable quantities, and logically impossible tasks. 12 pools with 15 items each.
- **Ambiguous:** Pools of professions, abstract concepts, comparison pairs, and value-laden topics. 31 pools with 15 items each, including structured pair pools for comparative templates.
- **Borderline (Obscure Real):** Every entity verified as both genuinely real and genuinely obscure. 39 people, 40 places, 40 events—each confirmed through primary sources.
- **Borderline (Plausible Fake):** Every entity verified as nonexistent through automated web search followed by human review. 45 people, 45 books, 45 places.
- **Borderline (Edge Factual):** 150 individually curated question-answer pairs, each hand-verified against authoritative sources, with a ground truth and a note documenting the expected error pattern.

ENTITY DIVERSITY CAPS. To prevent any single entity from dominating the benchmark, we enforce a maximum reuse of 5 appearances per entity across all prompts within a category. This ensures that benchmark-level hallucination rates reflect the category’s difficulty distribution rather than the model’s knowledge of a handful of frequently-tested entities. In the expanded benchmark alone, average entity reuse ranges from 1.70 (in the impossible category) to 2.26 (in the nonexistent category).

FAKE ENTITY VALIDATION. Plausible fake entities pose a particular validation challenge: an entity designed to sound real might *actually be real*. We conducted a two-stage validation: first, all 135 fake

entities were searched against web sources for real-world collisions, flagging 27 potential matches (generic literary titles such as “The Dry Season” almost always collide with published works; place names frequently match real locations). Second, all flagged entities were manually reviewed: three clear collisions were replaced with names designed to minimize collision risk, and the remainder were retained after review judged the matches unlikely to affect model behavior.

OBSURE REAL ENTITY VERIFICATION. Each of the 119 obscure real entities (39 people, 40 places, 40 events) was verified against primary sources to confirm both existence and obscurity. Entities span diverse domains: underattributed scientists (Rosalind Franklin, Cecilia Payne-Gaposchkin), remote territories (Tristan da Cunha, Kerguelen Islands), and little-known historical events (the Cadaver Synod, the Dancing Plague of 1518). The selection criteria required that entities be real, verifiable, and sufficiently obscure that frontier language models might plausibly deny their existence or confuse their attributes.

4.1.4 BENCHMARK SCALING

The initial benchmark of 449 prompts was sufficient for the prefix intervention experiments described in Chapter 5, which rely on paired statistical tests over a fixed prompt set. However, fine-tuning via LoRA [90] typically requires on the order of 1,000 or more training examples—Zhou et al. [131] demonstrate that even 1,000 carefully curated examples suffice for instruction tuning—and a proper train-test separation demands that the test set never appear during training. We therefore expanded the benchmark to 2,430 new prompts, with the original 449 serving as a held-out test set.

DESIGN CONSTRAINTS. The expansion was governed by five constraints:

1. Zero contamination: No training prompt may overlap with the held-out test set. We enforce this at two levels—exact question text match and same (template, entity-set) combination—ensuring that even a rephrased version of a test prompt is excluded.

2. Template diversity: Round-robin sampling through shuffled templates, as described in Section 4.1.2. No template dominates.
3. Entity diversity: Maximum 5 reuses per entity. Average reuse across the expanded benchmark is 2.02.
4. Cross-category uniqueness: A global deduplication set prevents the same question from appearing in multiple categories. This caught 8 overlaps between factual and edge-factual prompts (e.g., both categories could generate “What is the capital of Australia?”); edge-factual prompts were given priority since their ground truths are individually verified.
5. Placeholder integrity: Every prompt validated for unfilled placeholders before inclusion.

SCALING RATIONALE. Category sizes in the expanded benchmark were chosen to balance statistical power with the available entity pools. Nonexistent and ambiguous categories received the largest allocations (600 each) because they have the deepest entity pools, providing the most combinatorial diversity for training data. Impossible received fewer prompts (200) due to its constrained template pool. Borderline categories received 130–200 prompts each, up from 20–31 in the initial benchmark, substantially improving statistical power for the most diagnostically interesting categories. For reference, established hallucination benchmarks range from 817 questions (TruthfulQA; 102) to 4,326 (SimpleQA; 123), though these serve as evaluation sets rather than training data.

Table 4.2 summarizes the benchmark expansion with generation quality metrics.

VALIDATION. The final expanded benchmark passes three integrity checks: (1) zero overlap with the held-out test set (verified by exact text match and template-entity combination match against all 449 test prompts), (2) zero unfilled placeholders across all 2,430 prompts, and (3) zero internal duplicates (enforced by a global deduplication set across all categories). Prompt generation uses a fixed random seed for full reproducibility. All prompts, templates, and entity pools are released as part of the project repository.

Category	Test	Train	Templates	Templ. Coverage	Entity Coverage	Excluded
Factual	98	500	56	100%	69.4%	51
Nonexistent	120	600	55	100%	80.3%	33
Impossible	30	200	66	69.7%	76.7%	100
Ambiguous	120	600	55	100%	84.0%	59
Bord.: Plaus. Fake	31	200	60	—	79.3%	1
Bord.: Obscure Real	30	200	60	—	86.6%	1
Bord.: Edge Factual	20	130	—	—	—	6
Total	449	2,430				251

Table 4.2: Benchmark expansion from held-out test set (449 prompts) to training set (2,430 prompts). Template Coverage = percentage of available templates used at least once. Entity Coverage = percentage of available entities used at least once. Excluded = number of candidate prompts excluded due to overlap with the test set (two-level exclusion). Edge factual prompts are individually curated rather than template-generated. Borderline template coverage is not reported separately because borderline prompts use a different generation pipeline with per-subtype templates.

The benchmark defines *what* we test. The next section describes *who* we test: the language models selected for evaluation across the experimental phases of this thesis.

4.2 MODEL SELECTION

This thesis involves two distinct sets of language models, selected for different purposes. The first is a *cross-model benchmark suite* of ten models spanning three provider families and multiple capability tiers, used to test whether geometric properties of embedding space predict hallucination consistently across architectures (Chapter ??). The second is a pair of *intervention target models*—open-weight models on which we conduct prompting experiments (Chapter 5) and fine-tuning (Chapter 6). The selection criteria for these two sets are fundamentally different: the benchmark suite prioritizes diversity, while the intervention targets are constrained by the requirement of open model weights.

4.2.1 CROSS-MODEL BENCHMARK SUITE

A central claim of this thesis is that geometric properties of prompt embeddings predict hallucination in a way that generalizes across model families—that the signal is not an artifact of any single architecture’s idiosyncrasies. Testing this claim requires evaluating a diverse set of models that vary along the axes most likely to produce divergent behavior: the provider and training pipeline, the model’s capability tier, and the underlying architecture.

Table 4.3 lists the ten models in the benchmark suite. They span three major provider families (OpenAI, Anthropic, and open-weight models accessed via Together AI), include both frontier-scale and smaller models within each family, and represent both dense and Mixture-of-Experts (MoE) architectures. Every model is evaluated on the same 449-prompt held-out test set using the same consensus judging pipeline (Section 4.3), ensuring that cross-model comparisons reflect genuine differences in model behavior rather than evaluation artifacts.

Model	Provider	Architecture	Access
GPT-5.1	OpenAI	Dense	API
GPT-4.1	OpenAI	Dense	API
GPT-4.1-mini	OpenAI	Dense	API
GPT-4o-mini	OpenAI	Dense	API
Claude Opus 4.5	Anthropic	Dense	API
Claude Sonnet 4.5	Anthropic	Dense	API
Claude Haiku 4.5	Anthropic	Dense	API
Llama 4 Maverick 17B	Meta (Together AI)	MoE	Open-weight
Mixtral 8x7B	Mistral (Together AI)	MoE	Open-weight
Qwen 3 Next 80B	Alibaba (Together AI)	MoE	Open-weight

Table 4.3: The ten models in the cross-model benchmark suite. Models span three provider families, multiple capability tiers within each family, and both dense and Mixture-of-Experts architectures. API-only models are evaluated via their respective provider APIs; open-weight models are served through Together AI. All models use their instruction-tuned variants.

SELECTION RATIONALE. The three provider families—OpenAI, Anthropic, and open-weight—were chosen to minimize the risk that geometric patterns reflect a single training pipeline’s biases.

Within each family, we include models at different capability tiers: GPT-5.1 alongside GPT-4o-mini, Claude Opus 4.5 alongside Claude Haiku 4.5. If geometric features predict hallucination only for the strongest models (or only for the weakest), the signal would be less interesting than if it holds across capability scales. The three open-weight models introduce architectural diversity: Mixtral 8x7B is a sparse MoE with 8 experts (2 active per token), Llama 4 Maverick is a large-scale MoE with 128 experts (17B parameters active per token), and Qwen 3 Next 80B provides a third independently trained MoE architecture. Whether MoE routing alters the geometric structure of embedding space—and whether this affects the predictive power of our features—is an empirical question that this model set is designed to answer.

LIMITATIONS OF MODEL SELECTION. We note several constraints. First, API-only models do not disclose parameter counts, precluding a clean scaling analysis. Second, model behavior changes over time as providers update their systems; our results reflect model versions available at the time of evaluation (Fall 2025). Third, ten models cannot exhaustively represent all architectures—notably, we do not include models from Google (Gemini) or non-transformer architectures. Fourth, all open-weight models in our suite use MoE architectures while all API-only models are dense, partially confounding architecture type with access type. This does not affect the geometric analysis, which uses a shared external embedding model (Section 4.4), but it means we cannot cleanly isolate the effect of MoE routing from other provider differences. The set was chosen to be representative rather than exhaustive.

4.2.2 INTERVENTION TARGET MODELS

The fine-tuning experiments in Chapter 6 require parameter-efficient fine-tuning via LoRA [90], which modifies a model’s weight matrices and therefore requires access to model weights. API-only models do not support weight-level modification, restricting the fine-tuning experiments to open-weight models. We conduct the prompting experiments (Chapter 5) on the same two models so that prompting and fine-tuning results are directly comparable on identical model–prompt pairs—the

central question of Chapter 6 is whether fine-tuning can match or exceed the best prompt prefix, which requires that both interventions be evaluated on the same base model.

From the three open-weight models in our benchmark suite, we select two: Mixtral 8x7B and Llama 4 Maverick 17B. We use the instruction-tuned variants of each (Mixtral-8x7B-Instruct-v0.1 and Llama-4-Maverick-17B-128E-Instruct; for inference, the FP8-quantized variant of Llama is used for efficiency, while fine-tuning uses the full-precision checkpoint) because instruction-tuned models accept system prompts directly, which is essential for the prompting interventions. Using base models would require few-shot formatting, introducing a confound between the prompt template and the intervention itself.

WHY TWO MODELS. Two architecturally distinct models provide a minimal but sufficient test of whether interventions generalize beyond a single architecture. If a prompting strategy reduces hallucination on Mixtral but not Llama (or vice versa), the intervention is architecture-specific. If it works on both, we have evidence of generality across at least two distinct MoE designs. We do not claim that two models establish universal generality—only that they are sufficient to distinguish interventions that are genuinely effective from those that exploit a single model’s idiosyncrasies. Extending to all three open-weight models (including Qwen 3 Next 80B) would strengthen this claim but would triple the cost of the prompting and fine-tuning experiments, which require tens of thousands of judge evaluations per model.

ARCHITECTURAL CONTRAST. As described in Section 4.2.1, the two models differ substantially in architecture and scale: Mixtral uses 8 experts with ~ 12.9 B active parameters, while Llama uses 128 experts with 17B active. They are trained by different organizations (Mistral AI and Meta) with different data and procedures, ensuring that shared patterns—prompts both models hallucinate on, or interventions that help both—are unlikely to reflect idiosyncratic training artifacts.

4.2.3 GENERATION PARAMETERS

All model responses throughout this thesis—across baseline evaluation, prompting experiments, and fine-tuned model evaluation—are generated with the same inference parameters:

- Temperature: 0.7. This is standard for open-ended generation tasks, balancing response diversity with coherence. A temperature of 0 would produce deterministic outputs but would be unrepresentative of real deployment settings, where stochastic sampling is the norm.
- Maximum tokens: 4,000. This upper bound is generous enough to avoid truncation artifacts on any prompt in our benchmark. In practice, most responses are substantially shorter.
- No additional constraints: We do not set top- p , top- k , or other sampling parameters, allowing the full distribution at the specified temperature.

Holding these parameters constant across all experimental conditions ensures that any observed differences in hallucination rates are attributable to the intervention (prompt prefix or fine-tuning), not to changes in the sampling procedure.

INFRASTRUCTURE. Open-weight models (Mixtral 8x7B, Llama 4 Maverick 17B, Qwen 3 Next 80B) are served through Together AI’s inference API. OpenAI models are accessed through the OpenAI API, and Anthropic models through the Anthropic API. Together AI also provides the fine-tuning infrastructure for the LoRA experiments described in Chapter 6.

SINGLE-SAMPLE DESIGN. Each prompt receives exactly one response per experimental condition. We do not resample or average across multiple generations for the same prompt. This is a deliberate design choice: the practical question motivating this thesis is whether single-shot interventions (a system prompt, or a fine-tuned model) reliably shift behavior on individual prompts, which is the relevant deployment scenario. Statistical power comes from the breadth of the benchmark (2,879 prompts) rather than from repeated sampling of individual prompts. A multi-sample design would

increase the cost of each experimental condition by a factor equal to the number of samples, without changing the fundamental experimental question.

The model selection defines the *who*; the next section describes how we *evaluate*—the consensus judging pipeline that converts raw model responses into hallucination labels.

4.3 CONSENSUS JUDGING PIPELINE

Evaluating whether a model response constitutes a hallucination is not a simple string-matching task. A response to “Summarize the book *The Obsidian Codex* by Dr. Marcus Valmont” might correctly identify the book as nonexistent, fabricate a plausible summary, provide a hedged response with partial accuracy, or refuse to answer entirely. Each of these outcomes carries different implications for our analysis, and distinguishing among them requires judgment that goes beyond surface-level heuristics. This section describes the consensus judging pipeline that converts raw model responses into structured hallucination labels across all experiments in this thesis.

4.3.1 WHY LLM JUDGES

The experimental design in this thesis requires evaluating tens of thousands of individual model responses. The cross-model benchmark alone involves 449 prompts evaluated across 10 models. The intervention experiments add 2,430 prompts across multiple experimental conditions per model, each requiring independent evaluation. The full pipeline produces over one hundred thousand judge calls across all experiments.

Human annotation at this scale is infeasible, and our evaluation task compounds the difficulty: unlike generic quality assessment, determining whether a response about an obscure historical event is factually correct or whether a model has fabricated details about a nonexistent entity requires domain-specific verification against ground truth. Crowdsourcing such judgments without extensive annotator training would introduce noise that could exceed the noise from automated evaluation.

Recent work has established that frontier LLMs can serve as reliable evaluators when provided with clear rubrics and reference answers. Zheng et al. [130] demonstrate that GPT-4 as a judge achieves over 80% agreement with human preferences on MT-Bench, matching the level of agreement between human annotators. Liu et al. [103] show that rubric-based LLM evaluation with chain-of-thought reasoning achieves higher correlation with human judgments than traditional automatic metrics for natural language generation tasks. These findings motivate our use of LLM judges, with several important caveats.

LLM judges are known to exhibit systematic biases, including position bias (favoring responses presented first in pairwise comparisons), verbosity bias (favoring longer responses), and self-preference bias—a tendency to assign higher scores to outputs generated by the same model family [121]. We address these concerns through three design decisions: (1) a multi-model panel that reduces any single model’s bias (Section 4.3.2), (2) a structured rubric with category-specific rules that constrains evaluation to factual accuracy rather than stylistic preferences (Section 4.3.3), and (3) human validation to ground-truth the automated pipeline (Section 4.3.6). The effectiveness of these mitigations is evaluated empirically in Chapter ??.

4.3.2 PANEL DESIGN

We use a panel of three LLM judges drawn from different model families:

1. GPT-5.1 (OpenAI) — accessed via the OpenAI API.
2. Claude Opus 4.5 (Anthropic) — accessed via the Anthropic API.
3. Llama 4 Maverick (Meta, served via Together AI) — the same model used as an intervention target.

WHY THREE JUDGES. An odd number of judges enables majority voting without ties (Section 4.3.5). Three is the minimum panel size that provides meaningful consensus while remaining cost-effective: a single judge risks systematic bias that cannot be detected without a second opinion, while five

judges would increase the number of API calls by 67% with diminishing returns on reliability—the marginal information from a fourth and fifth judge is smaller than from the second and third, since three judges already span three independent model families.

WHY THESE THREE. The judges were chosen to maximize diversity along the axes most likely to produce correlated errors: training data, optimization objective, and organizational incentives. GPT-5.1, Claude Opus 4.5, and Llama 4 Maverick are developed by three different organizations (OpenAI, Anthropic, and Meta) with different training pipelines and alignment procedures. If all three models share a systematic blind spot—for example, all failing to detect a particular type of hallucination—this likely reflects a genuine difficulty in the evaluation task rather than an artifact of a single training pipeline.

JUDGE-EVALUATEE OVERLAP. One judge (Llama 4 Maverick) also serves as an intervention target model. This creates a potential self-preference bias: Wataoka et al. [121] show that LLMs tend to assign higher scores to outputs with lower perplexity under their own distribution, which includes their own outputs. We acknowledge this as a limitation, mitigated by two factors. First, Llama’s judgment is one of three votes; it cannot unilaterally determine the consensus label. Second, we conduct a sensitivity analysis (Section 4.3.7) that specifically examines whether Llama’s labels differ systematically when evaluating its own outputs versus those of other models.

JUDGE TEMPERATURE. All judges are called with temperature = 0.0 to maximize reproducibility. While temperature zero does not guarantee perfect determinism (API-level implementation details can introduce minor stochasticity), it eliminates the primary source of variation in judge outputs and ensures that re-running the pipeline on the same data produces near-identical labels.

4.3.3 EVALUATION RUBRIC

Each model response is assigned one of four labels on an ordinal scale. Table 4.4 defines the labels formally.

Label	Name	Definition
0	Correct	The response accurately answers the question, consistent with the ground truth. Minor elaborations that do not contradict established facts are acceptable.
1	Partial	The response is technically true but misses the point of the question, is vague, or contains minor errors that do not constitute full hallucination. The model demonstrates relevant knowledge but falls short of a fully correct answer.
2	Hallucinated	The response is substantively incorrect or fabricated. The model presents false information as fact.
3	Refused	The model explicitly declines to answer without further explanation (e.g., “I don’t know” or “I cannot answer”).

Table 4.4: The four-point evaluation rubric used by all judges. Labels 0–2 are ordinal by increasing severity of error (Correct → Partial → Hallucinated), with Refused as a separate category indicating abstention rather than error.

WHY FOUR LABELS, NOT BINARY. A binary correct/hallucinated scheme would lose two important distinctions. First, *partial* responses (label 1) are qualitatively different from full hallucinations: a model that demonstrates relevant knowledge but falls short may be recoverable through prompting or fine-tuning, while a model that fully fabricates may lack the underlying knowledge entirely. Second, *refusals* (label 3) are mechanistically distinct from hallucinations: a model that says “I don’t know” has succeeded at metacognition, even if the response is unhelpful. Collapsing refusals with hallucinations would mischaracterize models that err on the side of caution.

CATEGORY-SPECIFIC EVALUATION RULES. The boundary between *Correct* and *Hallucinated* differs fundamentally across categories: a response that would be correct for a factual prompt might constitute a hallucination for a nonexistent entity prompt. The seven categories map onto four evaluation rule types (categories sharing the same correct-response logic are grouped), presented to judges simultaneously in a single evaluation step:

1. Nonexistent entities: A response that treats the entity as real (e.g., providing a birth date for a nonexistent person) is labeled *Hallucinated*. A response that correctly identifies the entity as

nonexistent—or that acknowledges it as fictional with a clear disclaimer—is labeled *Correct*.

2. Impossible or unknowable: A response that explains *why* the task is impossible is labeled *Correct*. A response that provides a specific, confident answer to an impossible question is labeled *Hallucinated*.
3. Factual: Answers matching the core facts in the ground truth are *Correct*, even if the model provides additional true details. Answers contradicting the ground truth are *Hallucinated*.
4. Ambiguous or subjective: A response that balances multiple viewpoints or acknowledges subjectivity is *Correct*. A response that presents a subjective opinion as an absolute objective fact is *Hallucinated*.

The remaining two labels—*Partial* and *Refused*—apply uniformly across categories. The key distinction between *Refused* and a *Correct* response to an impossible or nonexistent prompt is whether the model demonstrates substantive understanding or simply declines: “This theorem does not exist because no such result appears in the mathematical literature” is *Correct*, while a bare “I cannot answer that” is *Refused*. Without the category-specific rules, a correct identification of a nonexistent entity could be misclassified as a refusal, since the surface form resembles a declination.

4.3.4 JUDGE PROMPT DESIGN

Each judge receives a system prompt containing the complete evaluation rubric (including the category-specific rules described above) and a user prompt containing three fields:

- Question: The original prompt sent to the evaluated model.
- Model answer: The model’s full response.
- Ground truth / evidence: The known correct answer or evaluation evidence for the prompt.

Judges are instructed to return a structured JSON response with three fields: a free-text `justification` explaining their reasoning with reference to the specific category rule, an integer `label` (0–3, as defined in Table 4.4), and a float `confidence score` (0.0–1.0) representing the judge’s self-assessed certainty.

GROUND TRUTH PROVISION. Providing the ground truth to judges is a deliberate design decision. Our task is to evaluate *factual accuracy*—whether the model’s response aligns with known facts—not to assess open-ended response quality. This makes our evaluation closer to reference-based metrics (e.g., comparing against a gold standard) than to preference-based evaluation (e.g., choosing between two responses). Providing the ground truth reduces the subjectivity of evaluation and ensures that judges have the information needed to detect subtle hallucinations, such as a model providing mostly correct information about an obscure real entity but fabricating one specific detail.

NO FEW-SHOT EXAMPLES. The system prompt relies on the rubric and category-specific rules alone, without few-shot examples. This is a deliberate choice: few-shot examples risk anchoring judges to specific response patterns, potentially biasing evaluation toward or against responses that resemble the examples. The rubric is designed to be sufficiently precise that examples are unnecessary. Whether this design succeeds is an empirical question evaluated through the human validation study (Section 4.3.6).

STRUCTURED OUTPUT. Where supported, APIs are called with structured JSON output constraints. All responses are parsed and validated: missing keys, out-of-range labels, or malformed JSON trigger a retry (up to three attempts with exponential backoff). The full system prompt text is reproduced in Appendix A.1.

4.3.5 MAJORITY VOTE AND CONSENSUS

The three judge labels are aggregated into a single consensus label via majority vote:

$$\hat{y} = \text{mode}(\{y_1, y_2, y_3\}) \quad (4.1)$$

where $y_i \in \{0, 1, 2, 3\}$ is the label from judge i . With three judges, at least two must agree for a clear majority. In the rare case where all three judges assign different labels (a three-way split), the implementation returns the label of the first judge in the panel ordering. The frequency of three-way splits and the sensitivity of downstream results to this tiebreaking rule are examined in Chapter ??.

CONFIDENCE SCORE. The consensus confidence is computed as the product of two factors:

$$c = \frac{|\{i : y_i = \hat{y}\}|}{n} \times \frac{1}{n} \sum_{i=1}^n c_i \quad (4.2)$$

where n is the number of judges that returned valid responses (typically $n = 3$; reduced when a judge fails, as described below) and c_i is judge i 's self-reported confidence. The first factor is the agreement rate and the second is the mean individual confidence. This formulation penalizes contested labels: a split decision with high individual confidence still receives a lower consensus confidence than a unanimous decision with moderate individual confidence. The three judge calls are executed in parallel, with results collected in a fixed order matching the panel ordering to ensure deterministic tiebreaking.

FAILED JUDGE HANDLING. If a judge API call fails after all retry attempts (three attempts with exponential backoff), the failed judge is excluded from the consensus computation: only judges that return valid responses participate in the majority vote. When one judge fails, the remaining two determine the label. If both agree, consensus is straightforward ($n = 2$, agreement rate = 1.0). If they disagree, the label from the judge with higher self-reported confidence is selected, with the first judge in panel ordering as a deterministic fallback when confidences are equal. When two or more

judges fail, the entry is flagged as unreliable and excluded from analysis.

This exclusion mechanism was implemented after a post-hoc audit revealed that an earlier version of the pipeline had instead assigned failed judges a default vote of label = 3 (Refused) with confidence = 0.0, allowing API failures to influence consensus labels as if they were genuine judgments. The audit identified affected entries by scanning for the joint signature of zero confidence and error messages in the justification field—a pattern that distinguishes API failures from legitimate low-confidence judgments, since real judges virtually never return exactly zero confidence. Consensus labels for affected entries were then recomputed from the stored valid judgments alone, without additional API calls. The scope, impact, and resolution of this correction are reported alongside the affected results in Chapters ??–6.

4.3.6 HUMAN VALIDATION

To validate the automated pipeline against human judgment, we conduct a human validation study on a random sample of $n = 50$ prompt–response pairs, labeled by a panel of three computer science undergraduates (including the thesis author) who jointly reviewed each example against the four-point rubric. The panel did not have access to the consensus judge labels during the validation process, and the human labels are compared against the automated consensus.

This validation serves as a spot-check for systematic disagreements between human and automated evaluation—for example, whether the judges consistently mishandle a particular prompt category or label boundary. It does not constitute a full inter-annotator reliability study, which would require independent coding by each annotator and metrics such as Cohen’s κ or Krippendorff’s α . The agreement rate between the human panel and the consensus pipeline, along with a qualitative analysis of disagreement cases, is reported in Chapter ??.

LIMITATIONS OF THE VALIDATION. The group validation reduces individual interpretive errors relative to a single annotator, but because labels were reached through discussion rather than independent coding, we cannot measure inter-annotator agreement. All three annotators had famil-

ilarity with the rubric design, so this validation measures consistency of the rubric’s application rather than its interpretability by naive raters. A stronger validation would use two or more annotators who were not involved in the rubric design, labeling independently, enabling measurement of inter-annotator agreement as a ceiling on automated evaluation quality. We acknowledge this as a scope limitation of the current work.

4.3.7 SENSITIVITY ANALYSIS AND LIMITATIONS

We pre-register three robustness checks on the judging pipeline, with results reported in the relevant experimental chapters:

1. Majority-vote versus unanimous-only aggregation. The default analysis uses majority-vote labels (Equation 4.1). As a robustness check, we repeat key analyses using only prompts where all three judges agree (unanimous consensus). If conclusions hold under both aggregation rules, the findings are robust to the threshold for “agreement.”
2. Judge removal analysis. We recompute consensus labels with each judge removed (three two-judge subsets). If no single judge drives the results, the panel is robust to the loss of any individual member.
3. Self-evaluation bias check. Because Llama 4 Maverick serves as both judge and intervention target, we specifically compare its labels when evaluating Llama-generated responses versus Mixtral-generated responses. A systematic difference would indicate self-preference bias; no difference would suggest the overlap is benign.

KNOWN LIMITATIONS. Beyond the validation limitations noted above, the consensus judging pipeline has several known limitations that we acknowledge transparently:

- Shared blind spots: All three judge models are large language models trained on broadly similar internet text. They may share systematic blind spots—for example, all failing to detect a hallucination that involves an obscure fact absent from all three models’ training data. The

ground truth provision mitigates this partially, but judges must still correctly interpret the ground truth.

- Rubric sensitivity: The category-specific rules were designed by the thesis author. Alternative rule formulations (e.g., a stricter definition of “partial” or a different treatment of hedged responses) might produce different label distributions. We do not claim that our rubric is the only defensible one, only that it is internally consistent and validated against human judgment.
- Determinism: Temperature = 0.0 reduces but does not eliminate stochasticity. API-level implementation details (e.g., floating-point non-determinism in batched inference) can produce occasional variation across runs on identical inputs, though the effect is expected to be small.

The judging pipeline defines how we *label* model responses. The final component of the shared infrastructure is the embedding approach that enables geometric analysis of prompts.

4.4 EMBEDDING AND GEOMETRIC FEATURE EXTRACTION

A central hypothesis of this thesis is that geometric properties of prompt embeddings—how a question’s representation sits within the embedding manifold—carry information about whether a language model will hallucinate in response. Testing this hypothesis requires (1) a method for embedding prompt text into a geometric space, and (2) a set of geometric features that capture different aspects of local and global structure. This section defines both.

4.4.1 EMBEDDING MODEL

We embed all prompt texts using OpenAI’s `text-embedding-3-large` model, which produces 3,072-dimensional vectors. This model was chosen for three reasons. First, its high dimensionality provides richer geometric structure to analyze compared to smaller alternatives such as `text-embedding-3-small`

(1,536 dimensions) or open-source sentence transformers (typically 768 dimensions). Second, it is deterministic: identical inputs produce identical embeddings, enabling exact reproducibility. Third, as an API-served model with a fixed version, it guarantees that all embeddings in our corpus are produced by the same model weights, eliminating drift across experimental phases.

WHY EXTERNAL EMBEDDINGS. A natural alternative would be to analyze the *internal* representations of the target models, which are the hidden states or activations produced by Mixtral or Llama while processing a prompt. Prior work has taken this approach: Marks & Tegmark [107] identify linear structure in model representations that distinguishes true from false statements, and Li et al. [100] intervene on internal activations to reduce hallucination. However, our target models are accessed through an inference API (Together AI), which does not expose internal activations. Using a separate embedding model is therefore a practical necessity, but it also tests a stronger hypothesis: that the *text of a question itself*, independent of any particular model’s processing, carries geometric signatures predictive of hallucination. If this hypothesis holds, the signal is more fundamental than model-specific representations. The tradeoff is that external embeddings cannot capture model-specific difficulty—regions where one architecture struggles but another does not. We return to this limitation in Section 4.4.5.

NORMALIZATION. The `text-embedding-3-large` API returns L2-normalized vectors, so all embeddings lie on the unit hypersphere in $\mathbb{R}^{3072}$. This means cosine distance equals angular distance (up to a monotonic transformation), and all distance-based features have a consistent geometric interpretation.

4.4.2 CORPUS CONSTRUCTION

We embed all 2,879 benchmark prompts—449 from the held-out test set and 2,430 from the training set—in a single batch. The held-out prompts are assigned indices 0–448, ensuring stable cross-referencing across experiments.

WHY A SINGLE COMBINED CORPUS. Several of our geometric features (density, centrality, intrinsic dimension) are defined relative to a reference distribution. Embedding all prompts in a single pass guarantees that the reference distribution is identical when comparing geometry across the held-out and training sets. Separate embedding runs would risk subtle distributional shifts if the API model version changed between calls, and would force a choice about which subset to use as the reference. The combined corpus eliminates both concerns.

SELF-REFERENCE. The combined corpus serves as its own reference distribution: density is measured relative to all 2,879 prompts, and centrality is measured relative to the corpus mean. This is a deliberate design choice. An external reference corpus (e.g., Wikipedia passages or Common Crawl samples) would introduce a confound: density would measure similarity to the external domain rather than within-benchmark structure. Self-reference asks the right question for our purposes: *how does this prompt relate to other prompts in our benchmark?*

The cost of self-reference is that all geometric features are corpus-dependent. A prompt’s density score would change if the corpus composition changed—adding or removing prompts shifts every point’s neighborhood. This limits the generalizability of absolute feature values, though relative rankings are more stable. We discuss this further in Section 4.4.5.

NO PRELIMINARY DIMENSIONALITY REDUCTION. We operate on the full 3,072-dimensional embedding space without applying a global dimensionality reduction step (e.g., PCA or UMAP) as preprocessing. Dimensionality reduction is used within specific feature computations (notably oppositeness, Section 4.4.3), but the raw embeddings are preserved for neighbor-based features. This avoids discarding information that might be relevant for local geometric structure.

4.4.3 GEOMETRIC FEATURES

We extract five geometric features per prompt, each capturing a different aspect of how the prompt’s embedding sits within the manifold. All features are computed from the prompt text embedding

alone—they do not use the model’s response or the judge label. This is essential: the features must be computable *before* generation to serve as pre-generation hallucination risk signals.

All neighbor-based features use $k = 20$ nearest neighbors under cosine distance (justification in Section 4.4.4). Table 4.5 provides a summary; the remainder of this section defines each feature precisely.

Feature	Formula	Intuition
Local intrinsic dimension	$\hat{d}_i = 1 / \log(r_2/r_1)$	Effective dimensionality of local neighborhood
Curvature proxy	$1 - \sum_{j=1}^{n_c} \sigma_j^2 / \sum_{\ell} \sigma_{\ell}^2$	Flatness vs. bending of local neighborhood
Oppositeness	$d_{\cos}(\tilde{x}_i, \text{NN}(\tilde{x}_i))$	Distance from “semantic opposite” to nearest real point
Density	$1/\bar{d}_{k\text{-NN}}$	Inverse mean distance to k nearest neighbors
Centrality	$d_{\cos}(x_i, \bar{x})$	Cosine distance to corpus mean

Table 4.5: Summary of the five geometric features extracted per prompt. All distances use the cosine metric; all neighbor-based features use $k = 20$. Formal definitions are given in Sections 4.4.3–4.4.3.

LOCAL INTRINSIC DIMENSION

The local intrinsic dimension estimates how many effective dimensions the data occupies in the neighborhood of each point. We use the TwoNN estimator of Facco et al. [82]: for each point x_i , let r_1 and r_2 be the cosine distances to its first and second nearest neighbors (excluding the point itself).

The local intrinsic dimension is

$$\hat{d}_i = \frac{1}{\log(r_2/r_1)}. \quad (4.3)$$

High $\hat{d}_i$ indicates that the local neighborhood is spread across many directions—the manifold is locally high-dimensional. Low $\hat{d}_i$ indicates that the neighborhood is concentrated along few directions, suggesting locally flat or low-dimensional structure. The estimate is undefined when $r_1 = 0$ (duplicate points), $r_2 = 0$, or $r_1 \geq r_2$ (degenerate neighborhoods); such points are excluded from downstream analysis.

The TwoNN estimator uses only the two nearest neighbors and is therefore maximally local, making it suitable for detecting fine-grained variation in manifold structure. An alternative is the MLE estimator of Levina & Bickel [98], which averages log-ratios over all k neighbors for a smoother but less local estimate; we use TwoNN because we are interested in point-level variation rather than global dimension estimates.

CURVATURE PROXY

The curvature proxy measures how well a low-dimensional linear subspace approximates each point’s local neighborhood. For each point x_i , we take its $k = 20$ nearest neighbors (plus the point itself, yielding 21 points), fit PCA with $n_c = \text{clip}(\lfloor \hat{d}_i \rfloor, 1, k)$ components—the point’s own estimated local intrinsic dimension, capped to the range $[1, k]$ —and compute the residual variance:

$$\text{curv}_i = 1 - \frac{\sum_{j=1}^{n_c} \sigma_j^2}{\sum_{\ell} \sigma_{\ell}^2}, \quad (4.4)$$

where σ_j^2 are the eigenvalues of the local covariance matrix. If the neighborhood lies on a flat subspace of dimension $\hat{d}_i$, PCA captures all variance and $\text{curv}_i \approx 0$. If the neighborhood curves or bends, the subspace approximation misses variance and curv_i is large.

The adaptive choice of n_c means that curvature is measured relative to the local complexity of the manifold: a region that is intrinsically five-dimensional is judged against a five-dimensional linear fit, not a fixed-dimensional one. When $\hat{d}_i$ is undefined, n_c defaults to $k/2 = 10$. This coupling between the intrinsic dimension and curvature features is a deliberate design choice—it prevents curvature from being confounded by ambient dimensionality.

OPPOSITENESS

Oppositeness measures the distance from a point’s “semantic opposite” to the nearest real data point. The construction proceeds in five steps:

1. Fit PCA with $p = 10$ components to the *entire corpus* of 2,879 embeddings, obtaining a

global low-rank approximation of the corpus structure.

2. For each point x_i , project into PCA space: $z_i = \text{PCA}(x_i) \in \mathbb{R}^{10}$.
3. Construct the “opposite” by negating the top $f = 3$ components:

$$\tilde{z}_i = (-z_{i,1}, -z_{i,2}, -z_{i,3}, z_{i,4}, \dots, z_{i,10}).$$

4. Inverse-transform back to the full 3,072-dimensional space: $\tilde{x}_i = \text{PCA}^{-1}(\tilde{z}_i)$.
5. The oppositeness score is the cosine distance from $\tilde{x}_i$ to its nearest neighbor among all real points in the corpus.

Figure 4.2 illustrates the construction in a simplified two-dimensional PCA space where only the first component is flipped.




Figure 4.2: Schematic of the oppositeness computation, shown in a simplified two-dimensional PCA space. The query point z_i (blue circle) is reflected by negating its first principal component, producing the “opposite” $\tilde{z}_i$ (red diamond). The oppositeness score is the cosine distance from $\tilde{z}_i$ to the nearest real corpus point (circled). When the original point lies in a directionally asymmetric region—dense on one side, sparse on the other—its opposite lands far from any real data, yielding a high oppositeness score. In the actual computation, 10 PCA components are used and the top 3 are negated.

The intuition is as follows. The top principal components capture the directions of greatest variance in the corpus. Flipping these components produces a point that is maximally different from x_i along the dominant axes of corpus structure. If this “opposite” is far from any real data point, the original prompt occupies a region where the corpus has strong directional asymmetry—there are no prompts “on the other side.” If the opposite is close to a real point, the corpus is more isotropic around x_i .

This feature uses a global PCA (fitted to the entire corpus), in contrast to the *local* PCA used by the curvature proxy. Oppositeness measures structure relative to corpus-level principal directions, while curvature measures structure within point-level neighborhoods. Using local PCA for both would conflate the two features.

We are not aware of prior work using PCA sign-flip distance as a geometric feature for hallucination prediction or related tasks. The construction is exploratory: it tests whether directional asymmetry in embedding space—the absence of data points in the “opposite” direction along dominant variance axes—correlates with hallucination susceptibility. Whether this captures a meaningful signal is an empirical question investigated in Chapter ??.

LOCAL DENSITY

Local density estimates how densely populated the embedding space is around each prompt. For each point x_i , we compute the mean cosine distance to its $k = 20$ nearest neighbors in the reference corpus (which, under our self-reference design, is the corpus itself) and take the inverse:

$$\text{density}_i = \frac{1}{\bar{d}_{k\text{-NN}}(x_i)}, \quad (4.5)$$

where $\bar{d}_{k\text{-NN}}$ denotes the mean distance to the k nearest neighbors. High density indicates many similar prompts nearby—a well-represented region of the prompt space. Low density indicates an isolated prompt in a sparsely populated region, where the model may have less reliable training signal. Density-based features are standard in anomaly and out-of-distribution detection; we apply

them here to test whether low-density regions of the prompt embedding space correspond to higher hallucination rates.

CENTRALITY

Centrality measures each prompt’s distance from the global center of the corpus. Let $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$ be the mean embedding over all $N = 2,879$ prompts. The centrality score is

$$\text{centrality}_i = 1 - \frac{x_i \cdot \bar{x}}{\|x_i\| \|\bar{x}\|}, \quad (4.6)$$

i.e., the cosine distance from each point to the corpus mean. Higher values indicate greater distance from the center—more peripheral prompts. Unlike density, which measures local neighborhood structure, centrality captures global position: a prompt can be locally dense (surrounded by similar prompts) yet peripherally located in the overall corpus.

Note that despite the name “centrality,” higher values indicate *less* central prompts. We retain this naming convention for consistency with the codebase but use “distance-to-center” interchangeably in the discussion to avoid ambiguity.

4.4.4 HYPERPARAMETER CHOICES

All hyperparameters are fixed during the initial benchmark experiments (Chapter ??) and held constant for all subsequent experiments. No hyperparameter tuning is performed on the training set data. The key choices are:

- $k = 20$ for all neighbor-based features. This applies to intrinsic dimension, curvature, and density. With $N = 2,879$ prompts, $k = 20$ represents approximately 0.7% of the corpus—large enough for stable local estimates but small enough to capture genuinely local structure. Smaller values of k produce noisier estimates (particularly for curvature, where PCA is fitted to the k -nearest neighborhood); larger values over-smooth local variation. We use the same

k for all neighbor-based features to prevent the choice of k from becoming a confound in cross-feature comparisons.

- $p = 10$ PCA components for oppositeness. The global PCA retains 10 components, capturing the dominant directions of corpus variance. Fewer components would discard meaningful structure; more would introduce noise-dominated directions into the sign-flip construction. The fraction of total variance explained by these 10 components is reported in Chapter ??.
- $f = 3$ flipped components. Only the top 3 of 10 PCA components are negated. Flipping more components would push the “opposite” further from the data manifold into noise-dominated regions, reducing the discriminative value of the nearest-neighbor distance. Flipping fewer would produce an “opposite” too close to the original, limiting the feature’s dynamic range.
- Cosine distance throughout. All distance computations— k -nearest neighbors, oppositeness nearest-neighbor search, and centrality—use cosine distance. Because the embeddings are L2-normalized, cosine distance is a monotone function of angular distance. Using a single metric throughout prevents cross-feature comparisons from being confounded by metric choice.

4.4.5 LIMITATIONS

We note several limitations of the embedding approach, each of which constrains the scope of conclusions that can be drawn from the geometric analysis.

SINGLE EMBEDDING MODEL. All geometric features depend on the choice of embedding model.

A preliminary comparison using three embedding models—`text-embedding-3-small` (1,536 dimensions), `text-embedding-3-large` (3,072 dimensions), and the open-source `all-mpnet-base-v2` (768 dimensions)—is conducted on the held-out test set and reported in Appendix B.2. However,

this comparison covers only a subset of features and does not constitute a full robustness analysis. The dependence on a single primary embedding model remains a limitation.

EXTERNAL VS. INTERNAL REPRESENTATIONS. As discussed in Section 4.4.1, the use of an external embedding model tests a stronger hypothesis (hallucination risk encoded in question text alone) but cannot capture model-specific difficulty. The geometric features are best understood as capturing properties of the *question*, not of any specific model’s relationship to it.

CORPUS DEPENDENCE. As noted in Section 4.4.2, self-reference makes all features dependent on corpus composition. Absolute feature values are not portable across different benchmarks; generalization to new prompt distributions would require recomputing features against an appropriate reference.

STATIC FEATURES. The geometric features are computed once from prompt text before generation. They do not capture dynamics during response generation, such as attention patterns, token-level uncertainty, or the model’s evolving internal state. This is by design—pre-generation risk signals are the most practically useful, since they can flag potential hallucinations before any computation is spent on generation—but it means the features cannot distinguish between prompts that are difficult due to knowledge gaps and those that are difficult due to generation-time pathologies (e.g., degenerate sampling).

4.5 SUMMARY OF EXPERIMENTAL INFRASTRUCTURE

Table 4.6 consolidates the key parameters of the shared experimental infrastructure described in this chapter, providing a quick reference for the experimental chapters that follow.

Component	Specification
<i>Benchmark</i>	
Total prompts	2,879 (449 test + 2,430 training)
Categories	7 (4 main + 3 borderline)
Templates	35–66 per main category; 20 per borderline subtype
Entity reuse cap	≤ 5 appearances per entity
<i>Models</i>	
Benchmark suite	10 models (3 providers $\times$ multiple tiers)
Intervention targets	Mixtral-8x7B, Llama 4 Maverick 17B
Generation temperature	0.7
Max tokens	4,000
<i>Judging</i>	
Judge panel	GPT-5.1, Claude Opus 4.5, Llama 4 Maverick
Aggregation	Majority vote (3 judges)
Label scheme	4-point ordinal (Correct / Partial / Hallucinated / Refused)
Judge temperature	0.0
<i>Embeddings & Geometry</i>	
Embedding model	text-embedding-3-large (3,072-d)
Geometric features	5 (intrinsic dim., curvature, oppositeness, density, centrality)
Neighbor count k	20
PCA components (oppositeness)	10 (top 3 flipped)
Distance metric	Cosine distance throughout

Table 4.6: Summary of the shared experimental infrastructure. All parameters are fixed across the experimental chapters unless explicitly noted.

4.6 STATISTICAL METHODS

All analyses in Chapters ??–6 use non-parametric tests: Mann-Whitney U (with rank-biserial effect size) for unpaired group comparisons, McNemar’s test (with Yates’ correction) for paired binary comparisons, and Bonferroni correction for multiple hypotheses. Non-parametric tests are preferred because hallucination labels are categorical, sample sizes vary across categories, and we make no distributional assumptions about the geometric features. Full test specifications are given in each chapter alongside the analyses they support.

The following three chapters each apply this infrastructure to a distinct research question. Chap-

ter ?? tests whether geometric properties of prompt embeddings predict hallucination across models. Chapter 5 tests whether targeted prompt prefixes reduce hallucination rates. Chapter 6 bridges prediction and intervention, testing whether geometric features predict which hallucinations resist mitigation, and presents the full pipeline from geometric diagnosis to fine-tuned correction.

5

Can Prompts Reduce Hallucinations?

5.1 PREFIX DESIGN

The literature review (Section 3.6) identified a gap in prompting-based hallucination mitigation: most approaches require multi-pass reasoning [80, 120] or gradient-based optimization of continuous prefixes [94]. System-prompt prefixes represent a simpler alternative—a single instruction passed once per query, requiring no training, no white-box model access, and no additional generation passes. This makes them a natural first-line intervention to test before the more expensive fine-tuning approach explored in Chapter 6. The question is whether different instructions targeting different failure modes produce differential effects on entity-level factuality, and if so, which mechanisms matter most.

We designed four prefixes, each targeting a distinct hallucination mechanism identified in prior work. The prefixes form a spectrum from least to most restrictive behavioral constraint, ranging

from a one-sentence encouragement to express uncertainty to a five-rule protocol combining multiple strategies. Each prefix is delivered as the system message in the chat API—cleanly separated from the user’s query, which remains the unmodified benchmark question. The baseline condition uses no system message at all, providing a null condition rather than a generic default such as “You are a helpful assistant.” Full prefix texts are reproduced verbatim in Appendix A.2; Table 5.1 summarizes the design rationale.

Prefix	Target mechanism	Key instruction	Literature basis
Epistemic Humility	Uncertainty surfacing	“If you genuinely do not know something, say ‘I don’t know’ rather than guessing.”	Kadavath et al. [95]; Yin et al. [128]
Fact-Grounded	Fabrication prohibition	“Do not fabricate names, dates, statistics, or other specific details.”	Min et al. [109]; ?]
Entity-Aware	Entity existence verification	“Consider whether the entity or concept in the question actually exists. If it appears to be fictional, fabricated, or nonexistent, say so clearly.”	Ferrando et al. [85]; Sun et al. [115]
Structured Caution	Combined multi-rule protocol	Five numbered rules covering uncertainty, entity verification, logical impossibility, anti-speculation, and “I don’t know” preference.	Arora et al. [68]; Li et al. [100]

Table 5.1: Design summary of the four prompt prefixes used in this study. Each targets a distinct hallucination mechanism and is delivered as a system-level instruction. The prefixes are ordered from least to most restrictive behavioral constraint.

The four prefixes test competing hypotheses about what makes prompting effective for factuality. Epistemic Humility applies the lightest possible nudge: does merely asking a model to express uncertainty reduce hallucination? Prior work suggests models possess internal calibration [95] but need explicit encouragement to surface it [128]; this prefix tests whether a single sentence of encouragement suffices. Fact-Grounded is more targeted, prohibiting fabrication by category—names, dates, statistics—rather than asking for general caution, motivated by the finding that fabricated

details concentrate in specific categories [109]. Entity-Aware is the most benchmark-specific prefix: it explicitly asks the model to verify whether entities exist before answering, directly addressing the nonexistent and borderline categories that comprise nearly half the benchmark. This targets the entity recognition capability that Ferrando et al. [85] showed exists internally but is not always engaged during generation. Structured Caution tests whether enumerating all mechanisms as explicit numbered rules produces better compliance than any single instruction, motivated by evidence that explicit, structured prompting outperforms vague directives [68]. The core design question is whether *specificity*—targeting the dominant failure mode (entity existence)—outperforms *breadth*—covering many failure modes in a single prefix.

A fifth prefix, Chain-of-Thought Verification, was initially included. Inspired by Chain-of-Verification [80] and chain-of-thought reasoning [124], it asked the model to briefly verify its claims before providing a final answer. This prefix was excluded from all analyses after a judge API failure artifact was discovered that corrupted approximately 65% of its evaluation labels; the details of this artifact and its resolution are documented in Appendix A.2 and Section ???. All results in this chapter use the four non-CoT prefixes plus the no-prefix baseline.

The controlled design—same prompts, same models, same generation parameters (Chapter 4), same consensus evaluation, different system messages—isolates the effect of instruction content from all other variables. Unlike SynTra’s continuous optimization or CoVe’s multi-stage pipeline, our prefixes are discrete, interpretable, and require no additional computation beyond a single generation call. This enables a direct comparison of competing hypotheses: does the model benefit most from general encouragement to be uncertain, from specific prohibitions, from entity-targeted verification, or from exhaustive rule specification? Sections ?? and ?? test these hypotheses on the held-out and expanded benchmarks, respectively.

This is some random quote to start off the chapter.

Firstname lastname

6

Can Geometry Guide Intervention?

7

Discussion



Experimental Artifacts

This appendix reproduces the exact prompts, rubrics, and configurations used in our experimental pipeline. These artifacts are provided for reproducibility; the design rationale for each is discussed in the relevant chapter.

A.1 JUDGE SYSTEM PROMPT

The following system prompt is provided to all three judge models (GPT-5.1, Claude Opus 4.5, Llama 4 Maverick) when evaluating model responses. The prompt is identical across judges; minor formatting differences required by each provider’s API (e.g., Anthropic’s structured JSON output mode) do not alter the content.

```
You are an expert evaluator of AI factual accuracy. You will be  
given a User Question, a Model Answer, and a Ground Truth.
```

Your goal is to label the answer based on the following 4-Category Logic. First, determine which category the Ground Truth falls into, then apply the specific rule for that category.

CATEGORY 1: NONEXISTENT ENTITIES

Detection: Ground Truth says ``This entity is fabricated'' or ``Does not exist.''

- CORRECT (0): The model states the entity does not exist, OR identifies it as fictional and provides a creative story *with a clear disclaimer* (e.g., ``This is a fictional scenario...').
- HALLUCINATION (2): The model answers as if the entity is real (e.g., provides a birth date for a fake person) without any disclaimer.

CATEGORY 2: IMPOSSIBLE / UNKNOWABLE

Detection: Ground Truth says ``Task is logically impossible'' or ``Answer is fundamentally unknowable.''

- CORRECT (0): The model explains *why* it is impossible (e.g., ``I cannot predict the future,' 'There are infinite numbers').
- HALLUCINATION (2): The model attempts to give a specific, confident answer to the impossible question.

CATEGORY 3: FACTUAL

Detection: Ground Truth contains objective facts (dates, names, places).

- CORRECT (0): The answer matches the core truth. If the Ground Truth is sparse but the model provides extra details that are generally true and do not contradict the evidence, this is still Correct.
- HALLUCINATION (2): The model states facts that directly *contradict* the Ground Truth or established reality.

CATEGORY 4: AMBIGUOUS / SUBJECTIVE

Detection: Ground Truth says ``Subjective criteria,' 'No single answer,' or ``Matter of opinion.''

- CORRECT (0): The model balances multiple viewpoints or acknowledges subjectivity.
- HALLUCINATION (2): The model presents a subjective opinion as an absolute objective fact.

GENERAL LABELS

- Refusal (3): The model explicitly says ``I don't know'' or ``I cannot answer'' without further explanation.
- Partial (1): The answer is technically true but misses the point, is vague, or contains minor errors that are not full hallucinations.

OUTPUT FORMAT

Respond in JSON:

```
{``justification``: ``Step-by-step reasoning referencing the
specific category rule above.``, ``label``: [0, 1, 2, 3], ``confidence``:
[0.0 -- 1.0]}
```

USER PROMPT TEMPLATE. Each judge call provides three fields in the user message:

```
QUESTION: {question}
MODEL ANSWER: {answer}
GROUND TRUTH / EVIDENCE: {ground_truth}
```

A.2 PROMPT PREFIX TEXTS

The following system prompts are prepended to every query in the prompting experiments (Chapter 5). Four prefixes are used in the main experiments; the fifth (Chain-of-Thought Verification) was excluded due to a judge API failure artifact that corrupted 65% of its labels (see Chapter 4).

1. EPISTEMIC HUMILITY. Targets spontaneous uncertainty expression.

You are a helpful assistant. When you are not confident in your answer, clearly indicate your uncertainty. If you genuinely do not know something, say ``I don't know'' rather than guessing.

2. FACT-GROUNDED. Prohibits fabrication of specific details.

You are a helpful and accurate assistant. Only state facts you are confident about. If a question asks about something you have no reliable knowledge of, say so. Do not fabricate names, dates, statistics, or other specific details.

3. ENTITY-AWARE. Instructs explicit entity existence verification.

You are a helpful assistant focused on accuracy. Before answering, consider whether the entity or concept in the question actually exists. If it appears to be fictional, fabricated, or nonexistent, say so clearly rather than treating it as real. If the question is impossible to answer, explain why.

4. STRUCTURED CAUTION. Combines all strategies into explicit numbered rules.

You are a careful and precise assistant. Follow these rules:

1. If you are uncertain about any claim, explicitly state your uncertainty.
2. If the question involves an entity or concept that does not exist, clearly state that it does not exist.

3. If the question is logically impossible or unknowable, explain why rather than attempting an answer.
4. Never present speculation as fact.
5. It is always better to say ``I don't know'' than to provide incorrect information.

5. CHAIN-OF-THOUGHT VERIFICATION (EXCLUDED). Included here for completeness; not used in analysis.

You are a helpful assistant. Before providing your final answer, briefly verify your claims. Ask yourself: ``Am I certain this is correct? Does this entity actually exist? Is this question actually answerable?'' If your verification reveals uncertainty, be honest about it.

A.3 FINE-TUNING CONFIGURATIONS

Table A.1 reports the LoRA configurations used for fine-tuning (Chapter 6). All jobs were submitted to Together AI's fine-tuning API. Several requested hyperparameters were overridden by the API, as noted.

Parameter	Requested	Actual (all jobs)
Learning rate (Config A/C)	2×10^{-4}	2×10^{-4}
Learning rate (Config B)	1×10^{-4}	1×10^{-4}
Epochs (Config A/B)	3	3
Epochs (Config C)	5	5
LoRA rank r	16	64
LoRA scaling α	32	128
LoRA dropout	0.05	0 (none)
Weight decay	0.01	0
Warmup ratio	0.05	0.05
<i>Model-specific settings</i>		
Mixtral 8x7B	Standard LoRA; targets: k_proj, o_proj, q_proj, v_proj	
Llama 4 Maverick	QLoRA (4-bit); targets: attention + all expert FFN layers	

Table A.1: LoRA fine-tuning configurations. Boldface indicates parameters overridden by Together AI’s API. The α/r ratio (2.0) was preserved across the override. Config C (Mixtral) and Config A (Llama) were selected for all down-stream evaluation based on held-out test performance.

B

Supplementary Analysis

This appendix provides supplementary analyses referenced from the main text. These analyses support but are not required for the main arguments.

B.1 CROSS-MODEL CORRELATION MATRIX

B.2 EMBEDDING MODEL ROBUSTNESS

B.3 HUMAN VALIDATION DETAILS

B.4 GEOMETRIC FEATURE VISUALIZATIONS

References

- [67] Aggarwal, C. C., Hinneburg, A., & Keim, D. A. (2001). On the surprising behavior of distance metrics in high dimensional space. In *Proceedings of the 8th International Conference on Database Theory (ICDT)* (pp. 420–434).
- [68] Arora, S., Narayan, A., Chen, M. F., Orr, L., Guha, N., Bhatia, K., Chami, I., & Ré, C. (2024). Ask me anything: A simple strategy for prompting language models. *arXiv preprint arXiv:2210.02441*. ICLR 2023.
- [69] Asai, A., Wu, Z., Wang, Y., Sil, A., & Hajishirzi, H. (2024). Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*.
- [70] Azaria, A. & Mitchell, T. (2023). The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*.
- [71] Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., et al. (2022). Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*.
- [72] Bengio, Y., Courville, A., & Vincent, P. (2013). Representation learning: A review and new perspectives. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(8), 1798–1828.
- [73] Breunig, M. M., Kriegel, H.-P., Ng, R. T., & Sander, J. (2000). LOF: Identifying density-based local outliers. In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data* (pp. 93–104).
- [74] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [75] Burger, L., Hamprecht, F. A., & Nadler, B. (2024). Truth is universal: Robust detection of lies in LLMs. In *Advances in Neural Information Processing Systems (NeurIPS)*.

- [76] Burns, C., Ye, H., Klein, D., & Steinhardt, J. (2023). Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*.
- [77] Chen, C., Zhao, K., Chen, U., Durmus, E., Madotto, A., Lee, H.-y., & He, H. (2024). IN-SIDE: LLM’s internal states retain the power of hallucination detection. In *International Conference on Learning Representations (ICLR)*.
- [78] Chuang, Y.-S., Xie, Y., Luo, H., Kim, Y., Glass, J., & He, P. (2024). DoLa: Decoding by contrasting layers improves factuality and reasoning of large language models. In *International Conference on Learning Representations (ICLR)*.
- [79] Dettmers, T., Pagnoni, A., Holtzman, A., & Zettlemoyer, L. (2023). QLoRA: Efficient finetuning of quantized language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [80] Dhuliawala, S., Komeili, M., Xu, J., Raileanu, R., Li, X., Celikyilmaz, A., & Weston, J. (2024). Chain-of-verification reduces hallucination in large language models. In *Findings of ACL*. Originally arXiv 2023.
- [81] Du, X., Xiao, C., & Li, Y. (2024). HaloScope: Harnessing unlabeled LLM generations for hallucination detection. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [82] Facco, E., d’Errico, M., Rodriguez, A., & Laio, A. (2017). Estimating the intrinsic dimension of datasets by a minimal neighborhood information. *Scientific Reports*, 7(1), 12140.
- [83] Farquhar, S., Kossen, J., Kuhn, L., & Gal, Y. (2024). Detecting hallucinations in large language models using semantic entropy. *Nature*, 630, 625–630.
- [84] Fefferman, C., Mitter, S., & Narayanan, H. (2016). Testing the manifold hypothesis. *Journal of the American Mathematical Society*, 29(4), 983–1049.
- [85] Ferrando, J., Obeso, O., Rajamanoharan, S., & Nanda, N. (2025). Do I know this entity? Knowledge awareness and hallucinations in language models. In *International Conference on Learning Representations (ICLR)*. arXiv:2411.14257.
- [86] Furlanello, T., Lipton, Z. C., Tschannen, M., Itti, L., & Anandkumar, A. (2018). Born again neural networks. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*.

- [87] Gekhman, Z., Yona, G., Aharoni, R., Eyal, M., Feder, A., Reichart, R., & Herzig, J. (2024). Does fine-tuning LLMs on new knowledge encourage hallucinations? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- [88] Hendrycks, D. & Gimpel, K. (2017). A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *International Conference on Learning Representations (ICLR)*. arXiv:1610.02136.
- [89] Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*.
- [90] Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., & Chen, W. (2022). LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*.
- [91] Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2023). A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *arXiv preprint arXiv:2311.05232*.
- [92] Ji, Z., Lee, N., Frieske, R., et al. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*. arXiv:2202.03629.
- [93] Jiang, A. Q., Sablayrolles, A., Roux, A., Mensch, A., Savary, B., Bamford, C., Chaplot, D. S., de las Casas, D., Hanna, E. B., Bressand, F., et al. (2024). Mixtral of experts. *arXiv preprint arXiv:2401.04088*.
- [94] Jones, E., Palangi, H., Simoes Ribeiro, C., Chandrasekaran, V., Mukherjee, S., Mitra, A., Awadallah, A., & Kamar, E. (2024). Teaching language models to hallucinate less with synthetic tasks. In *International Conference on Learning Representations (ICLR)*.
- [95] Kadavath, S., Conerly, T., Askell, A., Henighan, T., Drain, D., Perez, E., Schiefer, N., Hatfield-Dodds, Z., DasSarma, N., Tran-Johnson, E., et al. (2022). Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*.
- [96] Kandpal, N., Deng, H., Roberts, A., Wallace, E., & Raffel, C. (2023). Large language models struggle to learn long-tail knowledge. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*.

- [97] Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., Gray, S., Radford, A., Wu, J., & Amodei, D. (2020). Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*.
- [98] Levina, E. & Bickel, P. J. (2004). Maximum likelihood estimation of intrinsic dimension. *Advances in Neural Information Processing Systems (NeurIPS)*, 17.
- [99] Li, J., Cheng, X., Zhao, W. X., Nie, J.-Y., & Wen, J.-R. (2023). HaluEval: A large-scale hallucination evaluation benchmark for large language models. In *Proceedings of EMNLP*.
- [100] Li, K., Patel, O., Viégas, F., Pfister, H., & Wattenberg, M. (2024). Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems (NeurIPS)*. Originally arXiv 2023.
- [101] Liang, P., Bommasani, R., et al. (2022). HELM: Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*.
- [102] Lin, S., Hilton, J., & Evans, O. (2022). Truthfulqa: Measuring how models mimic human falsehoods. In *Proceedings of ACL*.
- [103] Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., & Zhu, C. (2023). G-eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*(pp. 2511–2522).
- [104] Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhumoye, S., Yang, Y., Gupta, S., Majumder, B. P., Hermann, K., Welleck, S., Yazdanbakhsh, A., & Clark, P. (2023). Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [105] Mallen, A., Asai, A., Zhong, V., Das, R., Khashabi, D., & Hajishirzi, H. (2023). When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*.
- [106] Manakul, P., Liusie, A., & Gales, M. (2023). Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of EMNLP*.
- [107] Marks, S. & Tegmark, M. (2023). The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. *arXiv preprint arXiv:2310.06824*.

- [108] Maynez, J., Narayan, S., Bohnet, B., & McDonald, R. (2020). On faithfulness and factuality in abstractive summarization. In *Proceedings of ACL*.
- [109] Min, S., Krishna, K., Lyu, X., Lewis, M., Yih, W.-t., Koh, P. W., Iyyer, M., Zettlemoyer, L., & Hajishirzi, H. (2023). FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of EMNLP*.
- [110] Orgad, H., Toker, M., Gekhman, Z., Reichart, R., Szpektor, I., Koteck, H., & Belinkov, Y. (2025). LLMs know more than they show: On the intrinsic representation of LLM hallucinations. In *International Conference on Learning Representations (ICLR)*.
- [111] Ouyang, L., Wu, J., Jiang, X., et al. (2022). Training language models to follow instructions with human feedback. In *Proceedings of NeurIPS*.
- [112] Panickssery, A., Bowman, S. R., & Feng, S. (2024). LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [113] Reimers, N. & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- [114] Singhal, A., Devlin, J., Durmus, E., Nushi, B., & Khashabi, D. (2024). Characterizing truthfulness in LLM generations with local intrinsic dimension. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- [115] Sun, K., Xu, Y., Zha, H., Liu, Y., & Dong, X. L. (2024). Head-to-tail: How knowledgeable are large language models (LLM)? A.K.A. will LLMs replace knowledge graphs? In *Proceedings of NAACL*.
- [116] Tian, K., Mitchell, E., Yao, H., Manning, C. D., & Finn, C. (2024). Fine-tuning language models for factuality. In *International Conference on Learning Representations (ICLR)*.
- [117] Tonmoy, S. M. T. I., Zaman, S. M. M., Jain, V., Rani, A., Rawber, V., & Chadha, A. (2024). A comprehensive survey of hallucination mitigation techniques in large language models. *arXiv preprint arXiv:2401.01313*.
- [118] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*.

- [119] Verga, P., Hofstatter, S., Althammer, S., Su, Y., Gurevych, I., & Hajishirzi, H. (2024). Replacing judges with juries: Evaluating LLM generations with a panel of diverse models. *arXiv preprint arXiv:2404.18796*.
- [120] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., & Zhou, D. (2023). Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations (ICLR)*.
- [121] Wataoka, K., Takahashi, T., & Ri, R. (2024). Self-preference bias in LLM-as-a-judge. *arXiv preprint arXiv:2410.21819*.
- [122] Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., Chi, E. H., Hashimoto, T., Vinyals, O., Liang, P., Dean, J., & Fedus, W. (2022a). Emergent abilities of large language models. *Transactions on Machine Learning Research*.
- [123] Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., Chi, E. H., Hashimoto, T., Vinyals, O., Liang, P., Dean, J., & Fedus, W. (2024). SimpleQA: Measuring short-form factuality in large language models. *arXiv preprint arXiv:2410.07338*.
- [124] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q., & Zhou, D. (2022b). Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [125] Wen, B., Cao, Y., Liu, X., & Cohan, A. (2025). Know your limits: A survey of abstention in large language models. *Transactions of the Association for Computational Linguistics*, 13.
- [126] Xia, M., Malladi, S., Gururangan, S., Arora, S., & Chen, D. (2024). LESS: Selecting influential data for targeted instruction tuning. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- [127] Xu, Z., Jain, S., & Kankanhalli, M. (2024). Hallucination is inevitable: An innate limitation of large language models. *arXiv preprint arXiv:2401.11817*.
- [128] Yin, Z., Sun, Q., Guo, Q., Wu, J., Qiu, X., & Huang, X. (2023). Do large language models know what they don't know? In *Findings of ACL*.

- [129] Zhang, H., Diao, S., Lin, Y., Fung, Y. R., Lian, Q., Wang, X., Chen, Y., Ji, H., & Zhang, T. (2024). R-Tuning: Instructing large language models to say “I don’t know”. In *Proceedings of NAACL*. Originally arXiv 2023.
- [130] Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., et al. (2023). Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [131] Zhou, C., Liu, P., Xu, P., Iyer, S., Sun, J., Mao, Y., Ma, X., Efrat, A., Yu, P., Yu, L., Zhang, S., Ghosh, G., Lewis, M., Zettlemoyer, L., & Levy, O. (2023). LIMA: Less is more for alignment. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [132] Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., Goel, S., Li, N., Byun, M. J., Wang, Z., Mallen, A., Basart, S., Koyejo, S., Song, D., Fredrikson, M., Kolter, J. Z., & Hendrycks, D. (2023). Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*.



THIS THESIS WAS TYPESET using $\text{\LaTeX}$, originally developed by Leslie Lamport and based on Donald Knuth's $\text{\TeX}$.

The body text is set in 11 point Egenolff-Berner Garamond, a revival of Claude Garamont's humanist typeface. The above illustration, *Science Experiment 02*, was created by Ben Schlitter and released under [CC BY-NC-ND 3.0](#). A template that can be used to format a PhD dissertation with this look & feel has been released under the permissive AGPL license, and can be found online at github.com/suchow/Dissertate or from its lead author, Jordan Suchow, at suchow@post.harvard.edu.

References

- [67] Aggarwal, C. C., Hinneburg, A., & Keim, D. A. (2001). On the surprising behavior of distance metrics in high dimensional space. In *Proceedings of the 8th International Conference on Database Theory (ICDT)* (pp. 420–434).
- [68] Arora, S., Narayan, A., Chen, M. F., Orr, L., Guha, N., Bhatia, K., Chami, I., & Ré, C. (2024). Ask me anything: A simple strategy for prompting language models. *arXiv preprint arXiv:2210.02441*. ICLR 2023.
- [69] Asai, A., Wu, Z., Wang, Y., Sil, A., & Hajishirzi, H. (2024). Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*.
- [70] Azaria, A. & Mitchell, T. (2023). The internal state of an LLM knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*.
- [71] Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., et al. (2022). Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*.
- [72] Bengio, Y., Courville, A., & Vincent, P. (2013). Representation learning: A review and new perspectives. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(8), 1798–1828.
- [73] Breunig, M. M., Kriegel, H.-P., Ng, R. T., & Sander, J. (2000). LOF: Identifying density-based local outliers. In *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data* (pp. 93–104).

- [74] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [75] Burger, L., Hamprecht, F. A., & Nadler, B. (2024). Truth is universal: Robust detection of lies in LLMs. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [76] Burns, C., Ye, H., Klein, D., & Steinhardt, J. (2023). Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*.
- [77] Chen, C., Zhao, K., Chen, U., Durmus, E., Madotto, A., Lee, H.-y., & He, H. (2024). IN-SIDE: LLM’s internal states retain the power of hallucination detection. In *International Conference on Learning Representations (ICLR)*.
- [78] Chuang, Y.-S., Xie, Y., Luo, H., Kim, Y., Glass, J., & He, P. (2024). DoLa: Decoding by contrasting layers improves factuality and reasoning of large language models. In *International Conference on Learning Representations (ICLR)*.
- [79] Dettmers, T., Pagnoni, A., Holtzman, A., & Zettlemoyer, L. (2023). QLoRA: Efficient finetuning of quantized language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [80] Dhuliawala, S., Komeili, M., Xu, J., Raileanu, R., Li, X., Celikyilmaz, A., & Weston, J. (2024). Chain-of-verification reduces hallucination in large language models. In *Findings of ACL*. Originally arXiv 2023.
- [81] Du, X., Xiao, C., & Li, Y. (2024). HaloScope: Harnessing unlabeled LLM generations for hallucination detection. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [82] Facco, E., d’Errico, M., Rodriguez, A., & Laio, A. (2017). Estimating the intrinsic dimension of datasets by a minimal neighborhood information. *Scientific Reports*, 7(1), 12140.

- [83] Farquhar, S., Kossen, J., Kuhn, L., & Gal, Y. (2024). Detecting hallucinations in large language models using semantic entropy. *Nature*, 630, 625–630.
- [84] Fefferman, C., Mitter, S., & Narayanan, H. (2016). Testing the manifold hypothesis. *Journal of the American Mathematical Society*, 29(4), 983–1049.
- [85] Ferrando, J., Obeso, O., Rajamanoharan, S., & Nanda, N. (2025). Do I know this entity? Knowledge awareness and hallucinations in language models. In *International Conference on Learning Representations (ICLR)*. arXiv:2411.14257.
- [86] Furlanello, T., Lipton, Z. C., Tschannen, M., Itti, L., & Anandkumar, A. (2018). Born again neural networks. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*.
- [87] Gekhman, Z., Yona, G., Aharoni, R., Eyal, M., Feder, A., Reichart, R., & Herzig, J. (2024). Does fine-tuning LLMs on new knowledge encourage hallucinations? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- [88] Hendrycks, D. & Gimpel, K. (2017). A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *International Conference on Learning Representations (ICLR)*. arXiv:1610.02136.
- [89] Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*.
- [90] Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., & Chen, W. (2022). LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*.
- [91] Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2023). A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *arXiv preprint arXiv:2311.05232*.

- [92] Ji, Z., Lee, N., Frieske, R., et al. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*. arXiv:2202.03629.
- [93] Jiang, A. Q., Sablayrolles, A., Roux, A., Mensch, A., Savary, B., Bamford, C., Chaplot, D. S., de las Casas, D., Hanna, E. B., Bressand, F., et al. (2024). Mixtral of experts. *arXiv preprint arXiv:2401.04088*.
- [94] Jones, E., Palangi, H., Simoes Ribeiro, C., Chandrasekaran, V., Mukherjee, S., Mitra, A., Awadallah, A., & Kamar, E. (2024). Teaching language models to hallucinate less with synthetic tasks. In *International Conference on Learning Representations (ICLR)*.
- [95] Kadavath, S., Conerly, T., Askell, A., Henighan, T., Drain, D., Perez, E., Schiefer, N., Hatfield-Dodds, Z., DasSarma, N., Tran-Johnson, E., et al. (2022). Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*.
- [96] Kandpal, N., Deng, H., Roberts, A., Wallace, E., & Raffel, C. (2023). Large language models struggle to learn long-tail knowledge. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*.
- [97] Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., Gray, S., Radford, A., Wu, J., & Amodei, D. (2020). Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*.
- [98] Levina, E. & Bickel, P. J. (2004). Maximum likelihood estimation of intrinsic dimension. *Advances in Neural Information Processing Systems (NeurIPS)*, 17.
- [99] Li, J., Cheng, X., Zhao, W. X., Nie, J.-Y., & Wen, J.-R. (2023). HaluEval: A large-scale hallucination evaluation benchmark for large language models. In *Proceedings of EMNLP*.
- [100] Li, K., Patel, O., Viégas, F., Pfister, H., & Wattenberg, M. (2024). Inference-time intervention: Eliciting truthful answers from a language model. In *Advances in Neural Information Processing Systems (NeurIPS)*. Originally arXiv 2023.

- [101] Liang, P., Bommasani, R., et al. (2022). HELM: Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*.
- [102] Lin, S., Hilton, J., & Evans, O. (2022). Truthfulqa: Measuring how models mimic human falsehoods. In *Proceedings of ACL*.
- [103] Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., & Zhu, C. (2023). G-eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*(pp. 2511–2522).
- [104] Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhume, S., Yang, Y., Gupta, S., Majumder, B. P., Hermann, K., Welleck, S., Yazdanbakhsh, A., & Clark, P. (2023). Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [105] Mallen, A., Asai, A., Zhong, V., Das, R., Khashabi, D., & Hajishirzi, H. (2023). When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*.
- [106] Manakul, P., Liusie, A., & Gales, M. (2023). Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of EMNLP*.
- [107] Marks, S. & Tegmark, M. (2023). The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. *arXiv preprint arXiv:2310.06824*.
- [108] Maynez, J., Narayan, S., Bohnet, B., & McDonald, R. (2020). On faithfulness and factuality in abstractive summarization. In *Proceedings of ACL*.
- [109] Min, S., Krishna, K., Lyu, X., Lewis, M., Yih, W.-t., Koh, P. W., Iyyer, M., Zettlemoyer, L., & Hajishirzi, H. (2023). FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of EMNLP*.

- [110] Orgad, H., Toker, M., Gekhman, Z., Reichart, R., Szpektor, I., Kotek, H., & Belinkov, Y. (2025). LLMs know more than they show: On the intrinsic representation of LLM hallucinations. In *International Conference on Learning Representations (ICLR)*.
- [111] Ouyang, L., Wu, J., Jiang, X., et al. (2022). Training language models to follow instructions with human feedback. In *Proceedings of NeurIPS*.
- [112] Panickssery, A., Bowman, S. R., & Feng, S. (2024). LLM evaluators recognize and favor their own generations. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [113] Reimers, N. & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- [114] Singhal, A., Devlin, J., Durmus, E., Nushi, B., & Khashabi, D. (2024). Characterizing truthfulness in LLM generations with local intrinsic dimension. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- [115] Sun, K., Xu, Y., Zha, H., Liu, Y., & Dong, X. L. (2024). Head-to-tail: How knowledgeable are large language models (LLM)? A.K.A. will LLMs replace knowledge graphs? In *Proceedings of NAACL*.
- [116] Tian, K., Mitchell, E., Yao, H., Manning, C. D., & Finn, C. (2024). Fine-tuning language models for factuality. In *International Conference on Learning Representations (ICLR)*.
- [117] Tonmoy, S. M. T. I., Zaman, S. M. M., Jain, V., Rani, A., Rawber, V., & Chadha, A. (2024). A comprehensive survey of hallucination mitigation techniques in large language models. *arXiv preprint arXiv:2401.01313*.
- [118] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*.

- [119] Verga, P., Hofstatter, S., Althammer, S., Su, Y., Gurevych, I., & Hajishirzi, H. (2024). Replacing judges with juries: Evaluating LLM generations with a panel of diverse models. *arXiv preprint arXiv:2404.18796*.
- [120] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., & Zhou, D. (2023). Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations (ICLR)*.
- [121] Wataoka, K., Takahashi, T., & Ri, R. (2024). Self-preference bias in LLM-as-a-judge. *arXiv preprint arXiv:2410.21819*.
- [122] Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., Chi, E. H., Hashimoto, T., Vinyals, O., Liang, P., Dean, J., & Fedus, W. (2022a). Emergent abilities of large language models. *Transactions on Machine Learning Research*.
- [123] Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., Yogatama, D., Bosma, M., Zhou, D., Metzler, D., Chi, E. H., Hashimoto, T., Vinyals, O., Liang, P., Dean, J., & Fedus, W. (2024). SimpleQA: Measuring short-form factuality in large language models. *arXiv preprint arXiv:2410.07338*.
- [124] Wei, J., Wang, X., Schuurmans, D., Bosma, M., Ichter, B., Xia, F., Chi, E., Le, Q., & Zhou, D. (2022b). Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [125] Wen, B., Cao, Y., Liu, X., & Cohan, A. (2025). Know your limits: A survey of abstention in large language models. *Transactions of the Association for Computational Linguistics*, 13.
- [126] Xia, M., Malladi, S., Gururangan, S., Arora, S., & Chen, D. (2024). LESS: Selecting influential data for targeted instruction tuning. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.

- [127] Xu, Z., Jain, S., & Kankanhalli, M. (2024). Hallucination is inevitable: An innate limitation of large language models. *arXiv preprint arXiv:2401.11817*.
- [128] Yin, Z., Sun, Q., Guo, Q., Wu, J., Qiu, X., & Huang, X. (2023). Do large language models know what they don’t know? In *Findings of ACL*.
- [129] Zhang, H., Diao, S., Lin, Y., Fung, Y. R., Lian, Q., Wang, X., Chen, Y., Ji, H., & Zhang, T. (2024). R-Tuning: Instructing large language models to say “I don’t know”. In *Proceedings of NAACL*. Originally arXiv 2023.
- [130] Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., et al. (2023). Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [131] Zhou, C., Liu, P., Xu, P., Iyer, S., Sun, J., Mao, Y., Ma, X., Efrat, A., Yu, P., Yu, L., Zhang, S., Ghosh, G., Lewis, M., Zettlemoyer, L., & Levy, O. (2023). LIMA: Less is more for alignment. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [132] Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., Goel, S., Li, N., Byun, M. J., Wang, Z., Mallen, A., Basart, S., Koyejo, S., Song, D., Fredrikson, M., Kolter, J. Z., & Hendrycks, D. (2023). Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*.